

Chapter 1: The Essentials



"You don't need to understand how the forklift works to run a move-in. But you do need to know where the freight targets are."

There is a particular kind of frustration that professionals experience with AI. They try it once. It gives a vague, generic answer. They think, *this thing is overhyped*. And they go back to doing things the slow way.

What they don't realize is that they handed over the keys to the most capable operator on the floor — and then gave that operator no floor plan, no target times, no exhibitor list, and no idea which hall they were standing in. Then they complained that nothing got built.

Everything in this chapter exists to prevent that experience. By the time you finish here, you will understand exactly why AI gives the answers it does, how to change those answers profoundly, and how to build a system that works for you automatically — even while you're on a red-eye between show sites.

We are going to cover seven ideas, and then one more that changes the ceiling entirely. Each one is a piece of equipment. Together, they constitute the complete operating system for working with AI at a professional level.

And this matters more at GES than it does at most companies. We produce **4,000+ live experiences a year**. We serve **150,000+ exhibitors**. We operate in **75+ countries** out of **24 global production and warehouse facilities**, with roughly **2,600 people** carrying the load. Every one of those events is a deadline that does not move. When you find a way to save twenty minutes on a task, you don't save twenty minutes — you save twenty minutes times a number so large it changes what the company is capable of.

The Brain: Understanding the Large Language Model


Let's start with the most important question: **What is an AI model, really?**

Strip away the marketing language and the science fiction associations. Here is what you need to know: a Large Language Model is, in its purest form, a **very large, very fast pattern-matching machine** trained on an almost incomprehensible amount of human text. Books, technical manuals, contracts, websites, codebases, operations documentation, safety standards, design theory, conversations — essentially, a significant portion of everything human beings have ever written down.

The result of all that training is something that functions, for practical purposes, like an extraordinarily high IQ mind that has read everything. Not memorized, exactly — it does not have a database of facts it looks up. It has *learned patterns* — the deep structural relationships between ideas, concepts, words, and reasoning steps across virtually every field of human knowledge.

Think of it this way: **the LLM is the brain. It is pure IQ.**

This is not a metaphor designed to make you comfortable. It is the most accurate way to understand what you are working with. When you ask Copilot a question, you are putting a question to something with genuine, broad intellectual capability — capability that rivals or exceeds the best-read person you have ever met, across almost every domain.



Don't take our word for it. **Try this:**

Go to lmarena.ai (<https://lmarena.ai>) — the AI Arena. This is a live, independent benchmark where the world's AI models are ranked by real users who evaluate their responses head-to-head, blind.

Look at the leaderboard. Notice the scores. Notice that models you've heard of — GPT, Claude, Gemini — are at the top. This isn't marketing. It is a crowd-sourced intelligence test conducted millions of times.

The takeaway: **the intelligence in these models is real, measurable, and remarkable.** Microsoft Copilot is built on this same class of frontier models. You have access to one of the highest-scoring intelligences in the world, woven into the tools you already use every day — the same Outlook you use to chase a late freight manifest, the same Excel you use to reconcile labor hours after move-out.

But here is the critical caveat that will define everything that follows: **raw IQ, without context, is useless.**

A brilliant person who has never seen a show floor cannot tell you whether your target times are realistic. They can tell you how targeting *works* in the abstract. They cannot tell you that your carpet crew is coming in behind a

late-arriving general session set and that hall C will be a bottleneck on Wednesday morning. That knowledge lives with you. Your job in this chapter is to learn how to hand it over.

Tokens: The Atoms of AI Language

Before we talk about context, we need to briefly understand how AI "reads" and "thinks." The unit of operation is not the word — it is the **token**.

A token is roughly the smallest chunk of meaning the AI processes. Words get split into tokens. "GES" is likely a single token. "Drayage" might be two: "dray" and "age." The word "exhibitor" might be a single token. "Antidisestablishmentarianism" might be six. Common short words are often single tokens; rare or long words get subdivided. Industry jargon — the vocabulary we live in — often splits in surprising ways.

Here is why this matters to you as a professional: **everything in AI — the cost of the request, the limits of what you can input, the speed of the response — is measured in tokens.** When IT says Copilot has a "context window" of a certain size, they are talking about tokens. When you upload an exhibitor service kit or a 90-page show operations manual to Copilot, it is converted to tokens before the AI reads it.

Go to platform.openai.com/tokenizer (<https://platform.openai.com/tokenizer>). Paste in a few sentences from your last organizer proposal, an exhibitor service kit, or an email you sent during move-in. Watch as the text gets highlighted in chunks — each colored chunk is a single token.

Now try our vocabulary: **drayage, marshaling yard, install and dismantle, General Service Contractor, carnet, Exhibitor Appointed Contractor**. Count the tokens on each. Notice how our industry's specialized language gets chopped into pieces the model has to reassemble.

You are now seeing AI's native language. This is how it actually reads your words — and it's a useful reminder that the model does not automatically know that "I&D" means install and dismantle at GES. Sometimes you have to say it plainly.

Understanding tokens is not just a curiosity. It directly connects to one of the most important concepts in this entire book.

Token Economics: The Cost of Thinking

Every time you send a message to Copilot, tokens go in (your prompt) and tokens come out (the response). In an enterprise deployment across a company our size — thousands of people across North America, Europe, the Middle East, and India — this happens at enormous scale, every hour of every day, in every time zone we operate in.


Token economics refers to the practical implications of this:

- **Longer inputs cost more** (in compute, and in some licensing models, in real dollars)
- **More focused prompts get better results** — because you're using your token "budget" efficiently
- **Large documents uploaded to Copilot** are chunked and tokenized before being read
- **Some capabilities are metered directly** — as you'll see later in this

chapter, Microsoft's newest agent layer, Copilot Cowork, is billed in usage-based **Copilot Credits**. Tokens stop being an abstraction the moment they show up on an invoice.

For you as a practitioner, the lesson is precision: a well-crafted 50-word prompt will almost always outperform a rambling 500-word prompt. You are not chatting with a friend who needs emotional context. You are allocating resources — the same way you'd allocate labor hours against a build schedule. Be specific. Be clear. The AI will do more with less when you speak precisely.

Think of it the way you think about freight. You don't ship air. You don't pay drayage on empty crate volume if you can avoid it. Same discipline applies here: pack the prompt tight, ship only what's needed, and get more value per pound.



Context Engineering: The Flashlight in the Dark Room

Now we arrive at the concept that will transform your relationship with AI more than any other. It is called **context engineering**, and it is the essential skill of the AI era.

Here is the most important analogy in this entire book. Read it slowly.

Imagine you have hired the most brilliant operations consultant in the world. This person has read every logistics textbook ever written, has advised the largest event producers on earth, understands union jurisdiction, customs documentation, structural engineering, and crowd flow. Flawless track record. You bring them to the venue. But when they arrive, you put them in a completely dark room in the back of the hall — no windows, no lights, no floor plan, no exhibitor list, nothing.

You then lean in through a slot in the door and ask: *"How should we sequence move-in for this show?"*

They can't answer. Not because they aren't brilliant. Because **they can't see anything**. They have all the IQ in the world, but zero information about your world. The result is a generic answer that could apply to any show in any hall in any country. Useless.

Now you hand them a flashlight. You shine it at a section of the room — and in that section, you have placed the floor plan, the freight target schedule, the advance warehouse receiving report, the labor forecast for aisle carpet and booth sets, the venue's dock configuration, the union jurisdiction rules for that city, and the organizer's general session timeline.

Now they can answer. And the answer they give is extraordinary — because it combines **world-class intelligence with specific, relevant context**.

The flashlight is the context. The AI is the IQ. Context engineering is the art of knowing what to put in the flashlight beam.

$\text{Valuable AI Output} = \text{World-Class IQ} + \text{Precise, Relevant Context}$

*Copilot gives you the IQ. **You provide the context.** This is not a limitation of the technology. It is the entire skill.*



This idea maps directly onto something you already do every day. When a new floor manager joins a show team mid-cycle, you don't hand them the job and walk away. You hand them the floor plan, the target sheet, the exhibitor list, the punch list from the last site visit, and the three things the organizer cares about most. That handoff *is* context engineering. You already know how to do this for humans. You are simply learning to do it for a new kind of teammate.

Context Rot: When Your Flashlight Gets Stale

There is a phenomenon called **context rot** that every serious AI practitioner needs to understand.

Imagine you are in a long conversation with Copilot — 40 exchanges deep, across the span of an hour. The early messages in that conversation were highly relevant and rich. But as the conversation grew, the AI had to start "forgetting" the earliest parts because they no longer fit in its active memory. Meanwhile, the conversation has meandered — you asked a tangential question about hotel block pickup, then got pulled into a graphics production question, then came back.

The result is that the context the AI is currently working with is a **degraded version of what you started with**. Critical early instructions have faded. The richness of your initial setup has been diluted by the noise of everything that

came after. You are getting less useful answers not because the AI got worse, but because its flashlight is now full of unimportant things and the useful documents have slipped out the back.

It's the same reason a show file gets messy by day four of move-in. Everything is technically still in there. Nobody can find anything.

The fix is simple: For important work sessions, start fresh. Provide your context at the top of a new conversation. Don't rely on continuity from a long prior exchange. Fresh context in, sharp output out.

The Context Window: Your AI's Working Memory

The **context window** is the total amount of information — measured in tokens — that the AI can hold in its "attention" at one time. Think of it as the size of the room the flashlight can illuminate. Modern frontier models have context windows ranging from 128,000 to over 1,000,000 tokens — enough to hold entire operations manuals, full exhibitor service kits, lengthy organizer contracts, or many hours of meeting transcripts.

Microsoft Copilot's context window for enterprise users is substantial and continues to expand. For practical purposes, you can upload full RFP responses, multi-page show specs, post-show reconciliation reports, and extensive email threads. The AI will read all of it.

But remember: **bigger context window ≠ better focus**. You want to give the AI the *right* context, not the *most* context. A precise flashlight beats a floodlight when you know what you're looking for. Dumping every document from a show's SharePoint library into a prompt is the equivalent of shipping the entire warehouse to the venue and sorting it on the dock. Technically possible. Operationally foolish.

The Persona: Your AI's Identity and Instructions

Here is a question: if you brought that genius consultant onto a show site, would you just set them loose with no briefing? Of course not. You would tell them: *"Here is who you are on this account. Here is your role. Here is how I want you to communicate with the organizer. Here are the constraints you must operate within."*

In AI, this is called the **persona** — or more technically, the **system prompt**. It is the foundational instruction set that defines how the AI behaves before you ask it a single question.

A well-crafted persona can transform your AI from a generic assistant into something that feels like a specialized expert built specifically for your role at GES.

In Microsoft 365 Copilot, you define your persona at:

[m365.cloud.microsoft \(https://m365.cloud.microsoft\)](https://m365.cloud.microsoft) → Settings → Personalization

This is where you tell your Copilot: who it is, what it knows about you, how it should respond, what tone to use, what it should prioritize. The AI will carry these instructions into every interaction.



Example persona for a show operations professional at GES:

You are a senior trade show operations expert with 20 years of experience as a General Service Contractor across North American **and** EMEA venues. You understand **move-in** and **move-out** sequencing, freight targets, drayage **and** material handling, advance warehouse **and** marshaling yard operations, union jurisdiction, **and** install & dismantle labor planning deeply. When I ask you questions, respond with the precision of a senior practitioner, **not** a generalist. Use clear, direct language. When I ask for analysis, give me a recommendation, **not** just a summary. Always flag safety, labor jurisdiction, **and** venue compliance considerations. My name is [Name] **and** I work in [Team] at GES.

Example persona for an exhibit designer or graphics production lead:

You are a senior environmental **and** exhibit designer working **inside** a global exhibitions **organization**. You think in terms of **brand** storytelling, attendee **journey**, sightlines, **and buildability**. You understand modular **and** custom stand construction, large-format graphics production, color accuracy across **substrates**, **and** sustainable **build** practices including reusable stand components. When I **share** a concept, push **back** on anything that will not survive contact with a real **show** floor, a real **budget**, **or** a real **install** crew. **Be** specific about materials **and dimensions**.

Example persona for a sales or account management lead:

You are a senior account **director** serving **show organizers**, exhibitors, **and** experiential agencies. You understand **multi-year organizer** partnerships, exhibitor satisfaction, sponsorship revenue strategy, **and** post-**show** reconciliation. Write in a warm, partnership-focused, operationally confident voice – never salesy, never hype-y. When drafting client communication, lead with what we will do **and by** when. Flag any commitment I am making that operations has not yet confirmed.

With a persona in place, every conversation starts with a fundamentally different AI than the generic one everyone else is using. You have built your own senior advisor.

***Trust** — we trust each other to always be honest and do what's right.*

A persona shapes how the AI talks. It does not make the AI correct. Trust in AI is the same as trust on a show site: earned through verification, not assumption. You would never ship a stand to a venue without checking the drawings against the space. Don't ship an AI output to an organizer without checking it either.

Meta-Prompting: Teaching Yourself Through AI

We need to talk about the most powerful skill in this entire book. It is called **meta-prompting**, and once you understand it, you will never interact with AI — or with information in general — the same way again.

Here is the insight: **AI doesn't just answer questions. It can simulate expertise you don't yet have.**

The standard approach to using AI is: *I have a question, I ask the AI, I get the answer.* That's useful. But it keeps you in the position of someone asking for answers.

Meta-prompting inverts the relationship. Instead of asking the AI what *you already want to know*, you ask it: *"From the perspective of an expert in X, what are the most important questions I should be asking about Y? What am I probably missing? What are the blind spots that non-experts in this field consistently have?"*

Let's make this concrete for GES professionals.

Open Copilot at [m365.cloud.microsoft \(https://m365.cloud.microsoft\)](https://m365.cloud.microsoft) and paste this prompt:


"From the perspective of a senior General Service Contractor operations executive, what are the five most important questions a show manager should ask before finalizing the move-in schedule for a 300,000 square foot exhibition in a venue we have not worked before? What blind spots do managers newer to that venue type typically have around freight targeting, dock capacity, and union jurisdiction? Be specific and direct."

*Read the response. Notice that it doesn't just give you answers — it reveals the **structure of expert thinking** in this domain. You just got a briefing from a senior operations mind without pulling a colleague off a live show to get it.*

Now try it with a topic from your own role. Replace the scenario with whatever is sitting on your desk right now — an international shipment needing a carnet, a sponsorship package that isn't selling, a graphics reprint that keeps missing color, a hotel block that isn't picking up.

This is meta-prompting. And its implications are staggering.

META-PROMPTING CONCEPT DIAGRAM: ITERATIVE IMPROVEMENT LOOP



The continuous cycle of using AI to optimize its own instructions for superior outcomes.

The deepest application of meta-prompting is **self-directed learning**. Any time you encounter a domain where you are not an expert — a new venue, a new country's customs regime, an unfamiliar vertical like aerospace or mining or medical devices, an emissions reporting standard, an emerging sponsorship model — you can use meta-prompting to rapidly acquire the *frame* of expert thinking in that domain.

This matters enormously at GES specifically. We operate in **75+ countries**. Nobody knows every venue, every union agreement, every customs rule, every cultural expectation. But everyone can now walk into an unfamiliar market with a working mental model in twenty minutes instead of three weeks.

You don't outsource your thinking. You **expand** it.

Understanding — *people come first; be understanding and compassionate.*

Meta-prompting is how you extend understanding across borders. Ask the AI what a first-time exhibitor is actually worried about. Ask what an organizer's board is under pressure to deliver. Ask what a local crew's constraints look like in a market you've never worked. Then go into the conversation as someone who already thought about the other person's day.

Your Voice Is Your Superpower: Wispr Flow and Super Whisper

Before we go further, we want to pause and introduce something that sounds simple but makes a profound difference in practice.

Change your relationship with AI. Start talking to it.

Right now, most people interact with AI by typing. They sit at a keyboard, laboriously craft a prompt word by word, and send it off. The result is often a shorter, less nuanced, less contextual prompt than what the person actually needed — because typing is slow and tedious and we abbreviate when we're tired.

For a workforce like ours, that's a real constraint. A lot of GES work does not happen at a desk. It happens on a show floor, in an aisle, at a marshaling yard gate, on a dock, in a hotel lobby at 6 a.m. before doors open. Typing a careful 200-word prompt while standing in hall D with a radio in one hand is not realistic. Talking is.

Wispr Flow ([wispr.flow \(https://wispr.flow\)](https://wispr.flow)) and **Super Whisper** ([superwhisper.app \(https://superwhisper.app\)](https://superwhisper.app)) are tools that let you speak to your computer's text fields — including the Copilot chat box — in natural speech. You think out loud, the tool transcribes, and your fully-formed thoughts appear as text prompts. No keyboard lag. No abbreviation. No loss of nuance.



People who adopt voice-first AI interaction report that the quality of their AI outputs increases dramatically. The reason is structural: humans speak at roughly 150 words per minute but type at only 40. When you speak, you provide more context, more nuance, more of the *why* behind your question — and the AI has so much more to work with.

*Download **Wispr Flow** (Mac/Windows) or **Super Whisper** (Mac). Set it up in 5 minutes. Then open Copilot, click in the chat box, activate your voice tool, and simply **talk** through a problem that's been sitting on your desk this week — the exhibitor escalation, the labor variance you can't explain, the design revision the client keeps circling.*

Don't overthink it. Just talk the way you'd explain the situation to a colleague standing next to you on the floor. Watch the prompt appear. Notice how much richer it is than what you would have typed. Send it.

This single habit change may be the highest-ROI thing you do this year.

Tools: Your Data Is Already Connected


Here is something that surprises most people when they learn it.

Copilot doesn't just have the LLM brain. **It already knows about your email. Your calendar. Your documents. Your Teams conversations. Your SharePoint files.** The moment you started using Microsoft 365 Copilot, all of those data sources were connected.

To understand why this matters, let's return to the Flashlight Theory. We said that the context is what makes the IQ valuable. The challenge, normally, is *getting* your context into the AI's flashlight beam. With most AI tools, you have to manually copy and paste your data, upload documents, and remind the AI who you are every time.

Microsoft solved this problem. They built the data connections — the technical bridges between the AI and your business data — directly into Copilot. In the AI world, these bridges are called **MCP Servers** (Model Context Protocol Servers). They are the technical standard for connecting tools and data sources to AI models.

You don't need to set any of this up. It is already done.



One of the most important things to understand about Copilot's data connections is the **permission inheritance principle**.

The AI can only access what **you** can access. If a document is restricted to another show team, it will not "bubble up" in your Copilot conversations. If an email thread is in someone else's inbox, Copilot won't surface it to you. Your existing Microsoft 365 permissions are inherited by the AI, automatically.

This is crucial for anyone handling client-confidential material at GES. Organizer contracts, exhibitor pricing, agency creative under NDA, attendee data governed by GDPR, unreleased show floor plans — your AI respects the same access controls that IT and Legal have configured. AI doesn't bypass security. It works within it.

Open Copilot at [m365.cloud.microsoft \(https://m365.cloud.microsoft\)](https://m365.cloud.microsoft). Now try each of these prompts, one at a time:

1. "Summarize the most important emails I received this week that require action from me before move-in."
2. "What meetings do I have coming up in the next three days, and what are the key topics?"
3. "Find any documents I've recently worked on related to [a show, organizer, or exhibitor name]."

Watch the AI pull from your actual data — your actual email, your actual calendar, your real files. This is not a demo. This is your show, viewed through AI.

Notice how different this feels from a generic AI assistant. You are not talking to a search engine. You are talking to something that knows your context — because your data is already in the flashlight beam.

Agents: Your First Synthetic Teammate

We have arrived at the most powerful idea in this chapter, and arguably in this entire book. It is the idea that will define competitive advantage in the live events industry for the next decade.

An agent is a synthetic teammate.

Not a chatbot. Not a search tool. A teammate — one that you configure, brief, and deploy to handle a repeatable set of tasks.

Here is the conceptual leap: we go from *prompts* (you ask a question, you get an answer) to *systems* (a configured AI that handles a class of work automatically, without you re-explaining everything every time).

A Copilot agent is built from four components:

{card} Persona

The system prompt — who is this agent, what is its role, how should it respond, what constraints must it operate within?

Files, documents, websites, meeting transcripts, and databases that the agent uses as its source of truth — exhibitor service kits, venue specs, safety standards, brand guidelines.

The connections it can make — can it search the web? Access SharePoint? Read a show document library? Create documents?

The operating procedure — how does it handle edge cases, escalations, and situations outside its knowledge?



To create an agent in Microsoft 365 Copilot, you go to **m365.cloud.microsoft**, click **New Agent**, and configure it:

1. **Name and description** — what this agent does
2. **Instructions** — the persona and operating procedures (this is your system prompt)
3. **Knowledge** — add files, websites, SharePoint links, meeting recordings, org charts
4. **Output capabilities** — enable it to create Word documents, Excel reports, PowerPoint presentations, or generate images
5. **Web access** — let it search the open web, or restrict it to only your specified sources

The agent you build will be available to you (and, if you choose, your team) as a named Copilot experience. Instead of crafting a new prompt every time you need the same type of work done, you simply open the agent and talk to it.

This is the moment we go from using AI to *deploying* AI.

We're Already Doing This

If this feels theoretical, it isn't. It's already happening inside the GES Collective.

In 2026, **onPeak** — our accommodations business — launched the **onPeak AI Smart Suite**: AI-powered hotel search, an AI contract reader that pulls terms out of hotel agreements, and intelligent email categorization that routes the flood of attendee and organizer correspondence to the right place. That is not a pilot deck. That is production software, built by our own people, running against real client work, under Lisa Baez's Tech-Enabled Services organization.

The same instinct shows up across the Collective: **GES Exhibitions** modernizing how we plan and run the floor, **Spiro** pushing custom design and production workflows forward, **Visit by GES** shipping next-generation NFC Touchpoints and data capture, **SHOWTECH** running power and lighting at scale, and **GES EMEA** building data-led emissions reporting on the road to net zero.

So the question is not *whether* GES does AI. We do. The question is whether *you* do — in your role, on your accounts, in your week.

Copilot Cowork: The Agent Layer You'll Actually Use

Building your own agent is powerful. But there is now something above it — a productized, enterprise-grade agent layer that Microsoft ships and supports directly.

On **June 16, 2026**, Microsoft made **Copilot Cowork** generally available. It is the most significant change to how work gets done inside Microsoft 365 since Copilot itself.

Here is what makes Cowork different from the Copilot chat experience you already know.

Regular Copilot is *conversational*. You ask, it answers. It drafts, you refine. It is fundamentally a very good assistant that hands things back to you.

Cowork is *executional*. You give it a complex, multi-step assignment, and it goes and completes it — across Outlook, Teams, Word, Excel, PowerPoint, and SharePoint — and returns **finished work, not drafts**.

The distinctions that matter:

Copilot Chat vs. Copilot Cowork



Dimension	Copilot (chat)	Copilot Cowork
What you get back	A draft, a summary, a suggestion	A completed deliverable
Task length	Seconds to a minute	Long-running — minutes to hours
Scope	One app, one ask at a time	Multi-tool, end-to-end across Microsoft 365
Where it runs	Interactive, with you present	Cloud-hosted — it keeps working when your laptop is closed
What it knows	Your prompt plus connected data	Grounded in Work IQ — your organization's work patterns, relationships, and content
Models	Frontier models	Multi-model , including Anthropic's Claude Opus 4.8 and Sonnet 4.6
Security	Microsoft 365 trust boundary	Same Microsoft 365 trust boundary — permissions, compliance, and data governance intact
Billing	Included with a Microsoft 365 Copilot license	Usage-based Copilot Credits , on top of the Copilot license

Read the "where it runs" row again, because it matters more to GES than to almost any other company.

Our people travel. Constantly. Show teams fly out on Saturday, work install through Tuesday, run the show Wednesday to Friday, and dismantle into the weekend. Laptops get closed and shoved in bags. Wi-Fi at a convention center is what it is. The idea that a serious piece of work — a post-show reconciliation, a 60-page RFP response, a multi-week labor variance analysis — can be assigned to a cloud-hosted agent that keeps working while you're in the air is not a productivity gimmick. It is a structural fit for how this industry actually operates.

Work IQ is Microsoft's layer of organizational understanding — who works with whom, which documents matter to which projects, how your teams actually communicate, what "the Q3 organizer review" refers to in your world.

For a company running 4,000+ events a year across 24 facilities, this is the difference between an agent that knows English and an agent that knows GES. Work IQ is what lets Cowork resolve "pull the recap deck for the show we just closed in Dubai" without you spelling out a SharePoint path.

Realistic Cowork assignments at GES might sound like:

- *"Go through the last six weeks of email and Teams messages on this account, build a post-show reconciliation summary in Excel comparing forecast labor hours to actuals by function, flag every variance over 10%, write a one-page narrative in Word explaining the drivers, and drop both in the show's SharePoint folder."*
- *"Read the organizer's RFP in SharePoint, pull our three most relevant past responses, draft a full response document in our voice, build the accompanying capabilities deck, and list every question we need the organizer to clarify before we submit."*
- *"Review every exhibitor service kit we published this quarter, identify inconsistencies in how material handling and drayage terms are explained, and produce a corrected standard language block plus a change log."*
- *"Analyze the last twelve months of freight cost per show by venue, build the pivot analysis, chart the outliers, and prepare a QBR deck section with three recommendations."*

Notice what all of these have in common: they are **real work**, they span **multiple applications**, they take **serious time**, and they end in a **deliverable someone can actually use**.

Two things to hold in your head at once.

First: Cowork runs inside the Microsoft 365 trust boundary. Same permissions, same compliance posture, same data governance. It sees what you see and nothing more.

Second — and this is the T.R.U.E. one: Responsibility means we deliver on our commitments. Cowork returns finished work. Finished is not the same as correct. If a Cowork-produced reconciliation goes to an organizer with your name on it, it is your reconciliation. Review it the way you'd review a stand before doors open — walk the whole thing, not just the front elevation.

And because Cowork is billed in usage-based **Copilot Credits** on top of your Microsoft 365 Copilot license, it is worth being deliberate about what you hand it. Use it for the heavy, multi-hour, multi-tool work where it earns its cost. Use plain Copilot chat for the quick stuff.

How to think about the layers:

- **Copilot chat** — your everyday thinking partner. Fast, conversational, free-form.
- **Custom agents** — your repeatable role-specific workflows, built by you.
- **Cowork** — your heavy-lift executor for complex, long-running, cross-application work.

You'll use all three. Most people start at the first, discover the second within a month, and reach for the third when they hit a task they genuinely didn't have time to do properly.

This exercise is the capstone of Chapter 1. You are going to build a working AI agent, configured for your specific role at GES. Budget 20 minutes.

Step 1: Define your agent's purpose. Think about a task you do repeatedly — preparing exhibitor briefing summaries, drafting organizer status updates, building move-in checklists for a venue,

reviewing design specs for buildability, summarizing post-show feedback, checking international shipping documentation requirements. Choose one.

Step 2: Open Copilot Studio at m365.cloud.microsoft (<https://m365.cloud.microsoft>) → **New Agent**.

Step 3: Write your persona. Use this template and customize it:

"You are a senior professional at GES, a global exhibition and experiential marketing services company founded in 1939 and headquartered in Las Vegas, producing 4,000+ live experiences a year for 150,000+ exhibitors across 75+ countries. You specialize in [your area]. When preparing [your task], you organize information clearly, flag safety and labor jurisdiction considerations, note anything that could affect move-in or move-out timing, and present recommendations in a format an organizer or account lead could act on immediately. Write in a warm, partnership-focused, operationally precise voice. Always be specific and direct. When you are uncertain, say so clearly rather than speculating."

Step 4: Add knowledge. Upload at least one document relevant to your work — an exhibitor service kit, a venue spec sheet, a show operations manual excerpt, a safety briefing, a design standards document, a brand guideline, or any reference file you routinely consult. If you don't have one handy, add the GES public website URL.

Step 5: Test it. Ask it a question you would normally research manually — something you'd otherwise ping a colleague about mid-show. Notice how different the response is compared to a generic Copilot conversation.

Step 6: Push one task to Cowork. Identify the single largest, most tedious, most multi-app task on your plate this month — the one you keep deferring because it needs three hours you don't have. Write it out as a complete assignment, with the source material and the desired deliverables named explicitly. That's your first real Cowork brief.

Step 7: Reflect. In your next team meeting, describe what you built and what it does. This is how AI transformation actually happens inside organizations — not from the top down, but from practitioners building tools that work.

Welcome to the next version of your role.

Why This Matters Right Now

There is a reason this book exists in this particular year and not five years ago.

On **December 31, 2024**, GES completed its separation from Viad Corp and became an independent company under Truelink Capital — **the first time in 55 years** that this business has controlled its own roadmap, its own investment decisions, and its own pace of change.

That is not a footnote. It is the whole context for everything in this book.

For five and a half decades, GES was a segment inside someone else's portfolio. Now we set our own priorities. We decide what to build, what to fund, what to modernize, and how fast. Companies rarely get a moment like this — a business with 85+ years of accumulated operational expertise, suddenly handed the keys to its own future, at exactly the moment a general-purpose technology arrives that multiplies what every single person in the company can do.

The tools in this chapter are how individuals participate in that. Not by waiting for a rollout. By learning the seven concepts, building an agent this week, and handing your first heavy task to Cowork.

Excellence — provide excellent service and execution.

Here's the honest framing. ***AI raises the floor. Humans raise the ceiling.***

AI will make the average first draft better, the average analysis faster, the average summary tighter. That's the floor. It will not stand in an aisle at 5 a.m. and make the judgment call that saves the show. It will not read the room in an organizer meeting. It will not notice that a client is stressed about something they haven't said out loud.

That's the ceiling. That's still us. AI just clears the busywork out of the way so more of your day is spent up there.

The Seven Concepts: Your Master Reference

Before we close Chapter 1, let's anchor everything you've just learned in a single summary you can return to:

Your AI Foundations — The Seven Essentials

Concept	What It Is	Why It Matters
The LLM (Brain)	A neural network trained on vast human knowledge — pure IQ	The engine behind Copilot's intelligence
Token	The atomic unit of AI processing (~¾ of a word)	Determines cost, speed, and context limits
Token Economics	The relationship between token usage and value/cost	Drives better, more precise prompting habits — and Cowork is billed this way
Context Engineering	The art of putting the right information in the AI's "flashlight"	The primary determinant of output quality
Context	The total tokens the AI can	Defines how much data you can

Window	hold in active attention	give Copilot at once
Persona / System Prompt	Foundational instructions that shape AI behavior	Turns a generic assistant into your specialized expert
Meta-Prompting	Using AI to reveal expert thinking and unknown unknowns	The highest-leverage cognitive skill of the AI era
Tools / MCP Servers	Data connections between the AI and your business systems	Already configured in Microsoft 365 — your data is live
Agents	Configured AI systems that handle repeatable work	Your first synthetic teammates — from prompts to systems
Copilot Cowork	Cloud-hosted, multi-tool agent that executes long-running work end to end	Returns completed deliverables, not drafts — the enterprise agent layer you'll actually use

CONCEPTS SUMMARY



Glossary

Large Language Model (LLM)

A neural network trained on massive text corpora that generates human-like text by predicting the most contextually appropriate next tokens. The cognitive engine behind Copilot.

Token

The atomic unit of text that an AI processes — approximately $\frac{3}{4}$ of an average English word. All AI costs, limits, and speeds are denominated in tokens.

Context Window

The maximum number of tokens an AI can process in a single interaction — its active "working memory." Modern frontier models support 128,000 to 1,000,000+ tokens.

Context Engineering

The practice of deliberately curating and structuring the information provided to an AI to maximize the quality and relevance of its outputs.

Context Rot

The degradation of context quality that occurs during long AI conversations as early, relevant instructions are displaced by newer, less relevant content.

Persona (System Prompt)

The foundational instruction set given to an AI that defines its role, tone, expertise, and behavioral constraints — set before any user messages.

Meta-Prompting

The practice of asking AI to reveal expert thinking frameworks, unknown unknowns, and structural knowledge about a domain, rather than simply answering a specific question.

MCP Server (Model Context Protocol)

The technical standard for connecting data sources and tools to AI models. Microsoft has pre-configured MCP connections between Copilot and all Microsoft 365 services.

Agent

A configured AI system combining a persona, knowledge base, tools, and instructions to handle a repeatable class of tasks — a synthetic teammate.

Copilot Cowork

Microsoft's cloud-hosted agent capability, generally available June 16, 2026. Executes complex, long-running, multi-tool tasks end to end across Outlook, Teams, Word, Excel, PowerPoint, and SharePoint, returning completed results rather than drafts. Multi-model — including Anthropic's Claude Opus 4.8 and Sonnet 4.6 — grounded in Work IQ, operating inside the Microsoft 365 trust boundary, and billed usage-based in Copilot Credits on top of a Microsoft 365 Copilot license.

Work IQ

Microsoft's organizational intelligence layer that grounds Copilot and Cowork in how your company actually works — people, relationships, projects, documents, and communication patterns.

Copilot Credits

The usage-based billing unit for Cowork consumption, charged in addition to a Microsoft 365 Copilot license.

Token Economics

The relationship between token consumption, cost, and value — the principle that precise, well-structured prompts outperform verbose, unfocused ones.

Copilot Studio

The Microsoft 365 interface for creating, configuring, and deploying custom AI agents within an organization's Copilot environment.

Permission Inheritance

The security principle by which Copilot agents respect existing Microsoft 365 access controls — the AI can only access data that the user themselves can access.

Wispr Flow

A voice dictation tool that transcribes spoken language into text in real time across any application, enabling voice-first AI interaction.

Super Whisper

A macOS voice transcription tool that converts speech to text across any text field, optimized for natural, conversational AI prompting.

GSC (General Service Contractor)

The official on-site service provider appointed by a show organizer to deliver exhibition services — floor plan execution, decorating, material handling, labor, and venue operations. GES's core role at thousands of events each year.

EAC (Exhibitor Appointed Contractor)

A contractor engaged directly by an exhibitor rather than by the show organizer, working on the show floor under the GSC's operational rules. Spiro frequently operates in this capacity.

Drayage / Material Handling

The movement of exhibitor freight from the dock or advance warehouse to the assigned booth space, and back out again at the close of the show.

Chapter Summary

You began this chapter knowing that AI is important. You end it knowing *why* — and more crucially, you know *how* to use it with precision and purpose.

The Large Language Model is pure IQ — extraordinarily capable, trained on the sum of human knowledge, and available to you through Microsoft 365 Copilot right now. But IQ without context is wasted. Context engineering — the art of putting the right information in the AI's flashlight beam — is the skill that determines everything. Your Microsoft 365 data is already connected to that flashlight. Your email, your calendar, your show documents — they are live context, already loaded.

Personas and meta-prompting multiply your capabilities, not by outsourcing your thinking but by extending the range of your expertise across 75+ countries and every vertical we serve. Agents take the system further: a synthetic teammate, configured by you, deployed for your work, available always. And Copilot Cowork takes it to its logical conclusion — an enterprise-grade executor that runs in the cloud, works while your laptop is closed, spans every app you use, and hands back finished work.

None of it replaces what makes GES what it is. **Trust, Responsibility, Understanding, Excellence** — those are human commitments, and they still belong to you. AI just clears the path so you have more room to keep them.

This is not the future. This is Tuesday morning during move-in.

In the next chapter, we go deeper into the mindset that makes all of this actually stick — and why the people who thrive in this transition aren't the most technical ones. They're the most curious.

Resources for Chapter 1

- *AI Arena (LLM Rankings):* lmarena.ai (<https://lmarena.ai>)
- *OpenAI Token Calculator:* platform.openai.com/tokenizer (<https://platform.openai.com/tokenizer>)
- *Microsoft 365 Copilot:* m365.cloud.microsoft (<https://m365.cloud.microsoft>)
- *Wispr Flow (Voice AI Input):* wispr.flow (<https://wispr.flow>)
- *Super Whisper (Mac):* superwhisper.app (<https://superwhisper.app>)
- *GES:* ges.com (<https://ges.com>)