

Chapter 5: Week 3, Session A — Introduction to Microsoft Copilot & Prompting Essentials



"The quality of your question determines the quality of your thinking."
— Peter Drucker

There is a version of Microsoft Copilot that is, frankly, underwhelming.

You type: *"Write me a report on trade show logistics."*

Copilot dutifully returns four paragraphs of polished, completely generic prose that could have been written about any exhibition, in any hall, in any year. It reads like a Wikipedia summary. It doesn't know what a marshaling yard is in practice, it has never watched a freight target slip by ninety minutes, and it has certainly never stood on an empty show floor at 4:00 a.m. wondering whether the aisle carpet is going to arrive before the exhibitors do.

You could have found better with a three-second web search. You close the window, return to your keyboard, and quietly conclude that AI is probably overhyped.

There is a second version of Microsoft Copilot — the version that the top performers in every organization that has deployed it eventually discover. In this version, you type a different kind of prompt. Copilot pulls from your actual emails, your recent meeting transcripts, the labor forecast your operations analyst shared yesterday, and the venue safety bulletin your show manager circulated last Friday. It gives you a first draft that reflects your tone, incorporates the actual freight weights from the exhibitor service kit, and accounts for the union jurisdiction rules that govern the hall you're working in. It saves you forty-five minutes. Then it saves you another forty-five. Then it rewrites the way you work.

The difference between those two outcomes is not the software. It is the prompt.

This chapter is about the architecture that makes the second outcome possible — how Microsoft 365 Copilot actually works, where to find it, how it connects to your organizational data — and then the five prompting techniques that separate average users from power users. Master these, and Copilot stops being a curiosity and starts being your most capable colleague.

*GES has been in the business of building experiences since **1939** — from a small sign and exhibit shop in Kansas City to a global operation headquartered at 7000 Lindell Road in Las Vegas. Today the GES Collective produces **4,000+ live experiences every year**, serves **150,000+ exhibitors**, operates in **75+ countries**, and runs **24 global production and warehouse facilities** with roughly **2,600 employees**.*

*Run the arithmetic. Four thousand events a year is more than **eleven shows per day**, every day, including the ones nobody wants to work. Every one of those shows generates floor plans, service kits, freight manifests, labor forecasts, safety briefings, change orders, and a post-show reconciliation. Every one of them has a hard, immovable date — the doors open whether you're ready or not.*

That is what makes prompting skill compound here in a way it doesn't compound at a company with soft deadlines. A forty-five-minute savings on a single document is a nice afternoon. A forty-five-minute savings replicated across 4,000 shows and thousands of colleagues is a structural change in how much a lean team can carry.

1. The Copilot Landscape: Not All Copilots Are Equal

Before we discuss prompting, we need to clear up a confusion that trips up most new users: **Microsoft has multiple products with "Copilot" in the name, and they are not the same thing.**

This is not a minor technical footnote. If you understand the distinction, you will know exactly which tool to reach for and why. If you don't, you'll end up frustrated that "Copilot" doesn't have access to your files — because you're using the wrong one.



Think of it this way: imagine a hospital. They have a public information desk at the front entrance — anyone can walk in, ask questions, and get general health information. That's useful. But when you need a specialist who actually knows your medical history, who can look at your chart and make recommendations based on *your* specific situation, you see a physician with access to your full record. That is a fundamentally different kind of help.

Copilot Chat (free, included in your Microsoft account) is the public information desk. It is web-connected, it can answer general questions, it can help you write a cover letter or explain a news story — but by default, it does not have access to your GES emails, your SharePoint show libraries, your Teams conversations, or any other organizational data. It is general-purpose. It is genuinely useful. But it is not your specialist.

Microsoft 365 Copilot (the paid license add-on) is the physician with your chart. It knows your organizational context. It connects to Microsoft Graph — the intelligence layer that maps your emails, your meetings, your chats, and your documents — and it uses that context to give you responses grounded in *your* actual work. When you ask it to summarize the key points from your call with a show organizer last Tuesday, it can actually do that. When you ask it to draft a move-in schedule based on the freight target spreadsheet your logistics coordinator shared this morning, it can pull that file.

Copilot Agents are a third category: specialized, purpose-built AI assistants configured to handle specific workflows. A Copilot Agent might be designed specifically to answer exhibitor questions from a show's service kit, or to automate an ESG emissions reporting workflow, or to triage inbound organizer requests against a standing scope of work. These are more advanced and, for most GES colleagues in this course, are the horizon toward which we're building — the payoff of foundational mastery.

If you want proof that this horizon is real and not theoretical, look inside your own company: the **onPeak AI Smart Suite** already ships AI-powered hotel search, an AI contract reader, and automated email categorization. GES is not deciding whether to adopt AI. GES is deciding how fast the rest of the organization catches up to the teams that already have.

{tab-item} Copilot Chat (Free)

What it is: A general-purpose AI assistant connected to the web.

What it can do:

- Answer general questions using web search
- Help draft documents, emails, and summaries
- Explain concepts, analyze text you paste in
- Work with content you explicitly share in the conversation

What it cannot do:

- Access your GES emails or calendars
- Pull from your Teams chats or SharePoint show libraries
- Reference organizational documents without you pasting them in

Best for: General research, personal productivity, drafting when you have content to share manually — for example, researching a venue's published rigging rules before a site survey.

Access: copilot.microsoft.com

What it is: An AI assistant grounded in your organizational context via Microsoft Graph.

What it can do:

- *Access your emails, meetings, files, and chats (only what you have permission to see)*
- *Draft documents referencing actual GES files in SharePoint/OneDrive — show operations manuals, exhibitor service kits, SOWs*
- *Summarize Teams meetings you attended, including cross-timezone show calls you slept through*
- *Create Excel analyses from data in your tables — labor hours forecast vs. actual, drayage weights, booth inventory*
- *Search across your entire M365 ecosystem*

What it cannot do:

- *Access data you don't have permission to view*
- *See other employees' private files or emails*
- *Go outside your Microsoft 365 security boundary*

Best for: Everything in this course. This is your primary tool.

Access: `m365.cloud.microsoft`

What they are: Purpose-built AI assistants configured for specific workflows or knowledge bases.

What they can do:

- Automate specific, repeatable workflows
- Answer questions from a defined knowledge base (e.g., a show's exhibitor service kit, or a venue-specific operations library)
- Handle multi-step processes without step-by-step prompting
- Integrate with specific systems beyond M365

What they require:

- Configuration by IT or a designated Copilot admin
- Clear definition of scope and data sources

Best for: Advanced use cases — exhibitor services Q&A, freight target triage, ESG emissions data collection, standardized post-show reconciliation packages.

Access: Through Microsoft Copilot Studio (admin-configured)

The Security Point That Should Give You Confidence

One of the most common concerns we hear from GES colleagues is: "If Copilot can see all my organizational data, can it share an organizer's confidential floor plan or an exhibitor's pricing with someone else?"

*The answer is no — and the architecture makes this structurally impossible. Microsoft 365 Copilot only surfaces information that the **signed-in user already has permission to access**. If an exhibitor services rep doesn't have permission to view another account team's client pricing files, Copilot cannot show them those files. The same access controls, conditional access policies, and compliance frameworks that govern your M365 environment govern Copilot. Your client data stays within the Microsoft 365 service boundary. It does not leave your tenant, it is not used to train Microsoft's AI models, and it is not accessible to other organizations.*

This matters enormously in a business where you may be simultaneously serving two organizers who compete for the same exhibitors, or where a Spiro custom build for one brand sits three aisles from a direct competitor's stand. Confidentiality isn't a nice-to-have at GES. It's the price of admission.

This is not marketing language. It is the architecture. Understanding it removes a legitimate concern and lets you focus on what Copilot can actually do.

"Trust each other to always be honest and do what's right."


Trust runs in two directions with AI. First, you should be able to trust the platform — and the permission architecture described above is why you can. Second, and more importantly, your colleagues have to be able to trust you when you hand them a Copilot-assisted document.

That means: never pass off an unverified draft as checked work. If Copilot generated the first version of a site survey writeup and you haven't walked the hall yet, say so. Honesty about the provenance of a document is not a weakness — it tells the next person how much scrutiny to apply. On a show floor where a wrong number becomes a wrong truck becomes a late move-in, that signal is worth a great deal.

2. The Front Door: m365.cloud.microsoft

Every powerful tool has an entry point. For Microsoft 365 Copilot, that entry point is **m365.cloud.microsoft** — the unified Microsoft 365 Copilot portal that brings together everything in one place.

Think of m365.cloud.microsoft the way you think of the service desk on a show floor: it's not the only way to get help (you can also call your account manager, use the online ordering portal, or grab the floor manager on the aisle), but it's the hub — the place where everything connects and where you can do the most, all in one location.



What you'll find at m365.cloud.microsoft:

The portal gives you access to the full Microsoft 365 Copilot experience from your browser — the same capability available through the Microsoft 365 Copilot desktop app and the mobile app. You can start conversations, manage ongoing threads, use the Prompt Gallery (more on this shortly), and access all Copilot integrations from a single authenticated session.

Three ways to access M365 Copilot:

- **Web:** m365.cloud.microsoft — works in any modern browser, no installation required
- **Desktop:** The Microsoft 365 Copilot app, available through your Microsoft 365 installation
- **Mobile:** The Microsoft 365 Copilot mobile app (iOS and Android) — the same Copilot, on your phone, with the same access to your organizational data

That third one deserves emphasis at GES more than at most companies. A large share of this workforce is not desk-bound. Show operations managers, labor coordinators, I&D leads, and freight coordinators live on their feet — in halls, in marshaling yards, at advance warehouses, in the back of a venue loading dock with bad Wi-Fi and a radio in one hand. **Mobile Copilot is not the consolation prize. For a significant part of GES, it is the primary interface.**

Bookmark it now. Open *m365.cloud.microsoft* in your browser and add it to your bookmarks bar. Then install the mobile app before your next show. You'll use it daily. It takes four minutes total, and it removes one more friction point between you and a tool that should feel like second nature by move-in day.

Copilot embedded in the apps you already use:

m365.cloud.microsoft is the standalone experience. But Microsoft 365 Copilot also lives inside every app in your M365 suite — and for most GES workflows, you'll use it there rather than switching to the portal.

- **In Word:** Copilot appears in the document margin. Ask it to draft a section, expand a paragraph, rewrite in a different tone, or summarize the document. You can reference a specific file from SharePoint by typing `/` and the file name in your prompt. Useful for RFP responses to organizers, show operations manuals, exhibitor service kits, SOWs, and safety briefings.
- **In Excel:** Copilot appears in the ribbon on the Home tab. It can analyze data in a table, generate formulas, create charts, identify trends, and surface outliers — without you needing to write a single formula manually. Think labor hour forecasts vs. actuals, material handling weights and drayage costs, booth inventory, show P&L, hotel block pickup rates.
- **In Outlook:** Copilot can draft email replies, summarize long email threads, and flag action items from your inbox — including the ninety-message exhibitor thread that spawned four change orders you haven't logged yet.
- **In Teams:** Copilot can summarize meetings you missed, recap decisions made during a call, and list action items — pulling from the meeting transcript in real time. For a company running teams across Las Vegas, Chicago, Toronto, London, Paris, Amsterdam, and Dubai, "summarize the call I couldn't take at 3:00 a.m. local" is not a party trick. It's a scheduling reprieve.
- **In PowerPoint:** Copilot can generate slide decks from a document, reorganize presentations, and suggest design improvements. Organizer pitch decks, post-show recap decks, sponsorship proposals, QBRs, design concept presentations.

The pattern across all of these is consistent: **you stay in the tool you're already using, and Copilot shows up as a natural part of the workflow.** You don't need to stop what you're doing and go somewhere else.

A Word About File Access in Word and Copilot

*Microsoft's official guidance is important here: for Copilot to access a file when you're prompting in Word, **that file must be stored in SharePoint or OneDrive** — not just on your local hard drive. If a file lives only on your desktop or a local folder, Copilot cannot reach it. This is by design — it's part of the security boundary.*

The practical implication for GES: store your working documents in SharePoint or OneDrive, not local drives. This is good practice regardless of AI — it ensures show documents are backed up, version-controlled, and accessible to the on-site crew who inherit them at move-in — but it's essential for Copilot to work as designed.

The failure mode here is extremely familiar to anyone who has worked a show: the definitive version of the floor plan lives on one person's laptop, that person is on a plane, and the hall is filling up. Cloud storage solves an operational problem first and an AI problem second.

Microsoft 365 Copilot Search:

The portal also includes **Microsoft 365 Copilot Search** — a universal search capability that works across all your M365 apps and connected third-party data. Instead of searching separately in Outlook, then SharePoint, then Teams, Copilot Search finds relevant content across your entire organizational ecosystem in a single query. For an account manager trying to find everything related to a specific organizer across emails, site survey notes, shared floor plans, prior-year reconciliations, and team chats — this alone is a significant capability.

Consider the concrete version. A show returns to the same venue every eighteen months. The institutional knowledge about that venue — the freight elevator that runs slow, the union jurisdiction quirk on the mezzanine, the dock that floods when it rains, the organizer's preference for carpet color — is

scattered across three years of email, two SharePoint sites, and the memory of a colleague who has since moved to another region. Copilot Search is how you recover that knowledge instead of rediscovering it the hard way.

3. Work IQ — The Intelligence Layer That Grounds Copilot in Your Context

Here is the feature that makes Microsoft 365 Copilot genuinely different from a general-purpose AI chatbot — and it is important enough that it deserves its own section.

When you ask a general AI tool a question, it answers from what it knows from training data — the internet, books, articles. That's valuable. But it doesn't know *you*. It doesn't know what happened in your pre-con call last Tuesday. It doesn't know the venue safety bulletin your operations team circulated last week. It doesn't know the revised booth dimensions the exhibit designer just shared via email.

Microsoft 365 Copilot uses **Microsoft Graph** to close that gap.

Microsoft Graph is the intelligence layer that maps your organizational context. Think of it as the connective tissue of your Microsoft 365 environment — it knows which emails you've sent and received, which meetings you've attended, which documents you've created and accessed, which chats you've had in Teams, and which files your colleagues have shared with you. It maps the relationships between all of this information and maintains that map in real time.

When you submit a prompt to Microsoft 365 Copilot, the system does something called **grounding**: before generating a response, Copilot accesses Microsoft Graph within your tenant to pull in relevant context from your actual work. If you ask "Summarize what I need to prepare for the pre-con call with the show organizer tomorrow," Copilot doesn't just give you generic meeting prep advice — it looks at the meeting invitation, pulls the agenda, finds relevant emails from recent weeks, locates the floor plan and labor estimate shared with the attendees, and synthesizes all of that into a personalized briefing.

This process — the prompt going in, the grounding against your Microsoft Graph data, and the response coming out — all happens within your organization's Microsoft 365 service boundary. Your data does not leave your tenant.

The Permission Principle: How Grounding Stays Safe

*The grounding process is governed by a principle that is worth understanding clearly: **Copilot only surfaces information the signed-in user already has permission to access.***

This is not a policy statement. It is structural. Copilot doesn't have a special administrative view of your organization's data that bypasses your existing access controls. It uses Microsoft Graph with the same permissions as your account. If you don't have access to a file in a show's SharePoint library, Copilot cannot include that file in your responses. If a colleague's calendar is set to private, Copilot cannot see the details of their appointments.

For GES, this means: the existing security model that governs your M365 environment — including Conditional Access, MFA, and all other compliance controls — governs Copilot's access. No new exposure. The same guardrails, applied to AI-assisted work. Client confidentiality obligations under your organizer and exhibitor agreements are not weakened by Copilot; they are enforced by the same permission structure you already rely on.

What grounding means in practice at GES:

- A **show operations manager** asks Copilot to draft the move-in sequencing narrative for an upcoming show. Copilot pulls the current floor plan revision from SharePoint, the freight target schedule shared by the logistics coordinator via email, and the notes from the last pre-con call — all automatically, because they're already in the user's M365 environment.
- A **logistics & freight coordinator** asks Copilot to flag any advance warehouse exceptions from the last 30 days of email. Copilot searches the user's inbox for relevant threads and surfaces a prioritized summary of late arrivals, weight discrepancies, and unlabeled shipments.
- An **account manager** asks Copilot to prepare talking points for an

organizer renewal conversation. Copilot finds the account history, prior-year post-show reconciliation, recent correspondence, and service scope documentation in SharePoint and synthesizes key points for the meeting.

- An **exhibit designer at Spiro** asks Copilot to summarize every client comment on a design concept across three review cycles. Copilot pulls the email threads and the Teams review meeting transcripts and returns a consolidated change log.
- An **onPeak housing analyst** asks Copilot to compare this show's room block pickup curve against the same event two years ago. Copilot locates both workbooks and surfaces the delta by hotel tier and booking window.

None of these require the user to manually attach files or paste content into the prompt. The intelligence layer does the retrieval automatically — because it already knows your work context.

The analogy that makes this click:

Imagine you have a brilliant new coordinator who has been working alongside you for three shows. They've sat in every pre-con call with you, read every email you've sent and received, reviewed every document in your shared show library, and walked the floor with you at every move-in. When you ask them to help you prepare for an organizer meeting, they don't need you to explain the history of the account — they already know it. They synthesize what they know into what you need.

That is what Microsoft Graph grounding does for Copilot. You don't need to brief it every time. It already knows your context. You just need to ask the right question.

The honest caveat: the analogy breaks down when it comes to judgment and craft. Your coordinator builds genuine understanding — they learn that this particular organizer says "flexible" but means "not really," and that the dock at this venue backs up on Tuesdays. Copilot's "knowledge" is a structured map of your data — incredibly useful, but not a substitute for the contextual judgment you bring. The output it generates is always a first draft, not a final verdict. Microsoft itself is explicit about this in its product documentation: *"Remember that Copilot generates a draft. You'll need to verify and modify details to make sure it's accurate and fits your tone and style."*

Read that sentence carefully. It is not a liability disclaimer buried in fine print. It is honest product guidance from Microsoft — and it is good professional practice.

"Be responsible for our actions and deliver on our commitments."

Grounding is powerful, and power invites a specific kind of laziness: accepting a well-formatted answer because it looks like it came from your files.

*It probably did. But "probably" is not a standard that survives move-in. If Copilot tells you the target time for Hall C freight is 06:00 and you put that in a document that 140 exhibitors will read, **you** own that number. Not Copilot. Not the person who shared the source spreadsheet. You.*


The practical habit: for any number, date, dimension, weight, or dollar figure that will drive an operational decision, trace it back to the source document before you ship the draft. Copilot makes the first ninety percent fast. Responsibility is the last ten percent, and it does not delegate.

4. The Microsoft 365 Copilot Prompt Gallery and Skilling Center

Before we get to the five prompting techniques, you need to know about two resources Microsoft has built specifically to help you get better at this faster.

AI MEMORY & PERSONALIZATION CONCEPT: CONTEXT ACCUMULATION

Stacked context layers growing over time, enhancing AI personalization.



The Copilot Prompt Gallery lives at m365.cloud.microsoft/copilot-prompts (<https://m365.cloud.microsoft/copilot-prompts>). It is a curated library of ready-to-use prompts organized by job function (finance, HR, operations, sales, legal) and by application (Word, Excel, Outlook, Teams). Every prompt in the gallery has been tested and verified — these are not examples someone invented in a slide deck. They are prompts that produce results in the actual M365 Copilot environment.

For GES colleagues, the gallery is your starting point, not your ceiling. Browse the operations, project management, and sales categories to find proven prompts you can adapt to your specific workflows. Then modify them using the five techniques in the next section to make them sharper, more specific, and more powerful. A generic "summarize this project status" prompt becomes far more valuable when you rewrite it as "summarize move-in readiness by hall, flagging any exhibitor whose freight has not hit the advance warehouse."

The Microsoft 365 Copilot Skilling Center is the official learning hub at adoption.microsoft.com/copilot/skilling-center (<https://adoption.microsoft.com/copilot/skilling-center/>). It includes structured learning paths, scenario-based guides, and role-specific content — all officially maintained by Microsoft. If you ever want to go deeper than this course takes you, the Skilling Center is the verified source.

How to Use the Prompt Gallery Right Now

- 1. Go to m365.cloud.microsoft/copilot-prompts*
- 2. Filter by your primary job function*
- 3. Find three prompts that address tasks you do regularly*
- 4. Rewrite each one in GES vocabulary — replace "project" with "show," "client deliverable" with "exhibitor service kit," "budget variance" with "labor forecast vs. actual"*
- 5. Try each one this week in the relevant M365 app*
- 6. Note what worked, what didn't, and what you'd adjust*

This is a fifteen-minute investment that will pay dividends in your first week of real Copilot use.

Build a GES prompt library while you're at it. Every time you land on a prompt that produces genuinely good output, paste it into a shared OneNote page or SharePoint list for your team. Within a month, a show operations team will have a house set of prompts covering pre-con prep, freight exception triage, safety briefing drafts, and post-show reconciliation summaries. That library becomes an onboarding asset — the fastest way to get a new coordinator to competence is to hand them the prompts your best people already use.


5. The Five Prompting Techniques

Now we get to the core of this chapter.

Microsoft 365 Copilot has access to your organizational context. It has sophisticated language models coordinating its responses. It is, architecturally speaking, a remarkable tool. But none of that matters if you don't know how to talk to it.

Prompting is not a technical skill. It is a communication skill. The same way a well-constructed question to a colleague gets you a better answer than a vague one, a well-constructed prompt gets you better output from Copilot. The techniques below are not tricks — they are frameworks for clarity. Learn them, practice them, and they become second nature.

If you have ever written a good set of I&D instructions, you already have the underlying instinct. A crew lead doesn't hand a team a note that says "build the booth." They specify the sequence, the tools, the tolerances, the finish standard, and what to do when something doesn't fit. Prompting is that same discipline applied to a different kind of collaborator.



Technique 1: Role-Based Prompting

The core idea: Before asking Copilot to do something, tell it who it is.

This sounds almost comically simple. It is also the technique that produces the most immediate and dramatic improvement in output quality for most users.

Here is why it works: language models like the one powering Microsoft 365 Copilot don't have a single fixed "voice" or perspective. They adapt their reasoning, vocabulary, tone, and framing based on the context they're given. When you assign a role — *"Act as a veteran general service contractor operations director with 20 years of show floor experience"* — you are not just changing the tone. You are shifting the entire conceptual frame from which Copilot approaches the task. The criteria it applies, the risks it looks for, the language it uses, the depth of analysis it attempts — all of these shift to match the assigned role.



The analogy: Think about what happens when you call a venue's main number with a question about rigging points. If you reach the general switchboard, you get a polite, general answer — maybe a PDF link. If you get routed to the venue's safety officer who signs off on rigging plans, you get a different kind of answer entirely — one that draws on specific expertise, that flags the load limits you didn't think to ask about, that speaks in the language of your actual problem.

Role-Based Prompting is how you route your request to the right expert, even when the expert is an AI.

Where the analogy breaks down: a real venue safety officer has genuine experiential judgment, knows the building's quirks, and carries professional liability for their sign-off. Copilot's role adoption is sophisticated pattern matching, not lived expertise. You still bring the judgment. Copilot brings the synthesis and first draft. **Never** substitute a role-played safety review for an actual one.

The revolution this enables: Consider what it means for an operations lead to prompt Copilot as a *"skeptical show organizer reviewing this service proposal line by line for anything that looks like padding."* Or for an exhibit designer to prompt as a *"marketing director at an exhibiting brand who has been told to cut booth spend by 15% and needs to justify every square foot."* Or for a logistics coordinator to prompt as a *"customs broker reviewing a carnet manifest for missing HS codes and value discrepancies."* Or for a

process improvement lead to prompt as a *"Six Sigma consultant looking for waste and bottlenecks in a move-out sequence."* Each role unlocks a different analytical lens — applied instantly, at scale, to your actual GES materials.

Prompting template:

Act **as** a [specific role **with** relevant background].
Your task **is to** [specific task].
[Additional context **about** what you need.]

GES examples:

Act as a cynical, veteran show operations director with 20 years of experience running general service contractor teams in convention centers across North America. Review the following **move-in** schedule **and** freight target plan **and** flag every assumption you would push back on in a pre-con call.

[Paste **or** reference the **move-in** schedule]

Act **as** a venue safety officer responsible **for** approving rigging **and** overhead hang plans. **Read** the **following** exhibit structural **summary** **and** identify **any language** that **is** ambiguous, **any load** figure that **is** unsupported, **and** anything that would cause you **to** withhold sign-off. Be specific about what documentation you would demand.

Act **as a** customs broker reviewing **an** ATA carnet manifest **for an** international show shipment. Read **the** following equipment list **and** flag every **item** that is likely **to** trigger **a** classification question, **a** valuation challenge, **or a** temporary-import restriction **at the** destination country.

Act as an account manager who has served this **show organizer** for five years. **Based** on the following account summary **and** recent correspondence, draft a renewal conversation agenda that acknowledges the relationship history **and** anticipates the **organizer's** likely concerns about material handling rates.

Act as an exhibitor services rep who fields 200 calls a day during **move-in** week. Read the following draft exhibitor service kit section **and** rewrite it so that the three questions you get asked most often are answered before they are asked.

Role-Based Prompting Power Move

Add a behavioral instruction alongside the role: "Be direct. Do not soften your feedback. Flag problems explicitly." Copilot's default tendency is to be somewhat diplomatic. In operations, you often need it to be blunt — the show floor does not reward diplomatic ambiguity about whether the truck is going to make the target. The role assignment plus the behavioral instruction together produce output that reads like a tough internal review, not a polished proposal.

"People come first. Be understanding and compassionate."

Role-Based Prompting is also the single best tool GES has for building empathy at scale — and empathy is the actual product here. An exhibitor who has flown in from another country, spent six figures on a stand, and has one shot to make their year deserves to be met by someone who understands their pressure.

Before you send a difficult email, prompt Copilot to "read this as the exhibitor receiving it on the morning of move-in, after their freight has already been delayed once." Before you finalize a policy, prompt as "a first-time exhibitor who has never read a service kit before." Before you publish an internal change, prompt as "an I&D crew member reading this on a phone at the end of a 14-hour shift."

Understanding isn't softness. It's accuracy about how your words will land on a real human being who is having a harder day than you are.

Technique 2: Chain-of-Thought Reasoning

The core idea: Ask Copilot to show its work before giving you the final answer.

This is one of the most counterintuitive techniques, because the instinct is always to ask for the answer — not the reasoning process. But in complex analytical tasks, asking for reasoning first produces dramatically better final answers. Here is why.

FREIGHT COST REASONING: A STRATEGIC CHAIN OF THOUGHT

Streamlining logistics decisions from origin to final quote.



When you ask Copilot directly for a conclusion, it pattern-matches toward the most statistically likely answer given the context. When you ask it to reason through the problem step-by-step, it builds each conclusion on the previous one — and that structured process tends to produce more coherent, internally consistent analysis. More importantly, it makes the reasoning visible, which means you can catch errors in the logic before they propagate into the conclusion.

The analogy: Think about a freight target quote versus the cost breakdown behind it. The quote — a single drayage number on a page — is the conclusion. The breakdown is the reasoning chain: inbound weight by shipment type, advance warehouse vs. direct-to-show split, special handling, overtime exposure based on the target window, marshaling yard dwell assumptions, and the contingency you built in for the two exhibitors who always ship late. When the exhibitor challenges the number, the quote alone gives you nothing to defend. The breakdown gives you a conversation.

The same holds for a labor estimate versus the forecast reasoning behind it. Anyone can produce a number of hours. What makes the number credible is the chain: booth count by size band, install complexity, hall configuration, union jurisdiction rules that dictate crew composition, historical productivity rates from the last time you did this show, and the weather contingency for the outdoor loading dock.

When Copilot shows its reasoning chain, you get the breakdown instead of just the number. That is almost always more useful, because you can engage with the analysis rather than just accepting or rejecting the conclusion.

Where the analogy breaks down: a real freight cost breakdown is built on actual measured weights and real rate cards, produced by people who are accountable for them. Copilot's reasoning chain is sophisticated generation — it should be reviewed for logical consistency, not treated as a substitute for verified operational data.

The prompting template:

Walk **me through** your reasoning step-**by**-step **before** giving **me the** final answer.
[Then state **the** task clearly.]

GES examples:

I need **to** evaluate whether **to** recommend that **an** exhibiting client ship direct-**to**-show **or** route through **the** advance warehouse **for an** upcoming event **with a** compressed **three**-day move-**in**.

Walk me through your reasoning step-**by**-step: what factors should I consider, what are **the** trade-offs **on cost, risk, and target timing**, **and** what questions **do** I need answered **before** making **a** recommendation?

Then give me your final recommendation framework.

A show organizer who represents roughly **\\$4.2M in** annual services revenue **is** asking us **to** match a competitor's proposal that includes lower material handling rates, expanded onsite staffing, **and a** dedicated exhibitor success team. We currently provide standard GSC scope **with a shared** services model.

Walk **me** through your reasoning **on** whether **to** negotiate, match, **or** hold our position — **then** give **me** a prioritized recommendation **with** the **key** risks clearly stated.

Our move-out **at a** recent show ran eleven hours past **the** contracted hall release **time**, **and** we absorbed **the** overtime. Walk me through your reasoning **on the likely contributing causes** – **sequencing**, carrier check-in, marshaling yard throughput, labor availability, exhibitor dismantle behavior – **and** only **then** give me **a** prioritized list **of the three** changes most likely **to** prevent **a repeat**.

At the model architecture level, Chain-of-Thought prompting works because it forces the language model to generate tokens that represent intermediate reasoning before generating the final answer token sequence. Each step in the chain conditions the next step — creating a more constrained, coherent generation path than a direct-to-conclusion prompt allows.

In plain language: when you ask for the answer directly, Copilot takes a shortcut. When you ask for the reasoning first, it has to build the bridge piece by piece, and those intermediate pieces keep it on track. The final answer that comes out of a reasoning chain is usually more defensible than one that appeared directly.


This is also why Chain-of-Thought is particularly valuable in contract- and compliance-adjacent GES work: the reasoning chain is auditable. You can show it to a senior colleague, an organizer's procurement lead, or an insurance reviewer and explain why the analysis landed where it did. In a business where post-show reconciliation disputes are settled by documentation, showing your work is not an academic virtue. It is how you get paid.

Technique 3: Few-Shot Prompting

The core idea: Show Copilot what "good" looks like before asking it to produce something new.

This technique is borrowed directly from machine learning — "few-shot learning" is the ability of a model to generalize from a small number of examples. Applied to prompting, it means you give Copilot two or three strong

examples of the output you want, and then ask it to produce a new output in the same style.



The analogy: Think about how you onboard a new coordinator who has strong general skills but no GES institutional knowledge. You don't hand them a style guide and tell them to write like the senior team. You hand them three or four examples of excellent post-show reconciliation reports written by your best people and say: *"Write at this level. Match this structure. Aim for this depth."* They look at the examples, internalize the pattern, and produce work that fits the house standard.

This is exactly what few-shot prompting does. The examples are your institutional standard, communicated directly to Copilot.

Where the analogy breaks down: your coordinator internalizes examples through genuine understanding, not pattern matching. Copilot extracts statistical patterns from your examples — which means if your examples have a systematic flaw or bias, Copilot will replicate that flaw. If every reconciliation report you feed it quietly buries the overtime variance in a footnote, every new one will too. The "garbage in, garbage out" principle applies here more directly than in most other techniques.

The prompting template:

Here are [number] examples of [type of output].
[Example 1 – labeled "Strong" or "Weak" as appropriate]
[Example 2]
[Example 3]

Now [state the task]. Match the style, depth, and structure of the strong examples.

GES examples:

Here are three examples **of** post-show reconciliation reports that our team considers high quality, **and** one example **of** a weaker report.

[EXAMPLE 1 – **Strong**: paste text]
[EXAMPLE 2 – **Strong**: paste text]
[EXAMPLE 3 – **Weak**: paste text]

Now **write** a post-show reconciliation report **for** the following **event** **using** the structure **and** depth **of** the **strong** examples. Here **is** the raw labor, freight, **and** change-order data:
[paste the data]

Here are two examples of the way our design team writes a booth design spec document – concise, dimensionally precise, **with** clear material callouts, no ambiguous qualifiers, **and** an **explicit** note **on** what **is** rental versus custom fabrication.

[EXAMPLE 1]
[EXAMPLE 2]

Now write a design spec document **for** the following concept:
[describe the concept]

Here are two examples **of** site survey writeups **from** venues we have worked **before**. **Notice** that **both** cover dock **configuration**, freight elevator capacity, ceiling **and** rigging constraints, **union** jurisdiction notes, aisle **and** egress requirements, **and** a "known surprises" section.

[EXAMPLE 1]

[EXAMPLE 2]

Now draft a site survey writeup **template for** a venue we have never worked, **and** list the questions I need **to** answer **on-site to** complete it.

Few-Shot Prompting and Confidential Data

When using examples to train Copilot's style, be thoughtful about what you paste in. Real organizer names, exhibitor pricing, negotiated rate cards, and non-public contract terms should be anonymized in your examples before pasting them into a Copilot prompt. Use placeholder names (Organizer A, Exhibitor X, Venue 1) and rounded figures in your examples.

This matters acutely at GES because you frequently serve competitors within the same vertical. The style instruction works just as well with anonymized content, and it keeps your prompting practice in alignment with GES's client confidentiality commitments and the T.R.U.E. value of Trust.

Advanced application — the "Four Examples" method:

For complex output like organizer proposals, sponsorship decks, or reconciliation packages, try four examples: two strong, one acceptable, one weak. Ask Copilot to explain what makes the strong examples better before generating new output. This forces a brief analysis step (a variation of Chain-of-Thought) that tends to sharpen the generation significantly.


There is a bonus here that teams underestimate: the explanation Copilot produces about *why* the strong examples are strong is frequently a better articulation of your house standard than anything currently written down. Capture it. You just generated the first draft of a style guide as a side effect.

Technique 4: Reverse Prompting

The core idea: Instead of you asking Copilot a question, ask Copilot to ask *you* questions.

This is the technique that most surprises people when they first try it — because it inverts the natural instinct. We assume we know what we need and we try to describe it. The problem is that we often don't know what we need precisely enough to describe it well. The resulting prompt is vague, the output is generic, and we conclude that Copilot isn't that useful.

Reverse Prompting short-circuits this by letting Copilot do the requirement-gathering.



The analogy: Think about the difference between an exhibitor walking up to the service desk with a vague request ("I need my booth to look better") versus an exhibitor services rep running a structured discovery ("What's the booth size and configuration? Is this an inline or an island? What's your target dwell time? Do you need power drops, and where? Is your freight already at

the advance warehouse or is it coming direct? What's your dismantle deadline?"). The rep who asks questions produces a dramatically better outcome than the one who takes the vague description and starts ordering furniture. Reverse Prompting makes Copilot the skilled services rep.

Where the analogy breaks down: a skilled rep's questions come from deep craft knowledge and an understanding of *your specific situation and venue*. Copilot's questions are good but may not always surface the most important uncertainty for your specific task. After it asks its questions, you should also add anything it didn't think to ask — the venue quirk, the organizer's unwritten preference, the crew constraint.

The prompting template:

I need your help **with** [general description **of** task].
Before you **start**, ask me **any** questions you need **to** perfectly understand my request **and** produce exactly what I need. Don't **proceed until I've** answered.

GES examples:

I need to write a presentation for our quarterly **business** review with the **show organizer's** leadership team. **Before** you start, ask me any questions you need to understand exactly what this presentation **should** accomplish, who will **be** in the room, what **show** performance data I have available, **and** what format works **best** for this audience.

I want **to** draft a proposal letter **to** a **show** organizer **we've been courting for six months in the healthcare association space**. **Before** writing anything, ask me everything you need to know about the relationship history, their current general service contractor arrangement, what scope **we're** proposing, **and** the tone I want.

I need **to** build an exhibitor communications plan **for** a **first-time** international show **where** roughly **40%** **of** exhibitors are shipping under carnet **from** outside **the** country. Before you draft anything, interview **me about** timeline, languages, customs constraints, target windows, **and** what our exhibitor services team can realistically support.

The Reverse Prompting Power Move

After Copilot asks its questions and you answer them, add one more line: "Is there anything else you need to know, or any assumption you're making that I should verify?" This second-level check often surfaces a critical variable that the first round of questions missed — particularly on complex, multi-stakeholder tasks where the organizer, the venue, the union, the EAC, and the exhibitor all have a say.

When Reverse Prompting works best:

- Complex, multi-part deliverables (organizer proposals, show operations manuals, sponsorship packages, strategy memos)
- Tasks where the audience or venue context matters significantly to the output
- Situations where you're not sure exactly what you want — but you know it when you see it
- Tasks you're doing for the first time with Copilot and haven't yet developed a strong prompt template for
- Any first-time venue, first-time organizer, or first-time country


Technique 5: The Sparring Partner

The core idea: Ask Copilot to push back on your ideas, not agree with them.

This is the most advanced technique in the set — and for many event and exhibition professionals, the most valuable. The default behavior of any AI assistant is to be helpful, which in practice means it tends to be agreeable. It

will draft your proposal, it will refine your service plan, it will polish your arguments. What it will not do, unless you specifically ask it to, is tell you where your thinking is weak.

The Sparring Partner technique flips this. You give Copilot a role defined by skepticism, and then you present your best thinking and ask it to attack it.



The analogy: Think about a full pre-con walkthrough with a hostile clock. You walk the floor plan with the venue, the organizer, and the union steward, and every one of them is looking for the thing you missed — the egress width that's four inches short, the hang point that isn't rated, the target window that assumes a dock can turn twelve trucks an hour when it has never done better than eight. The goal of that walkthrough is not to be right. It is to discover every weakness before move-in day, when the stakes are real and the cost of a fix multiplies by ten.

For GES professionals, the Sparring Partner technique is how you pressure-test your thinking before an organizer procurement review, an exhibitor escalation call, a leadership presentation, or a venue safety sign-off. Copilot, assigned the role of your toughest critic, will find the holes that your own confirmation bias is inclined to skip over.

Where the analogy breaks down: a real pre-con walkthrough involves people who know the building, know the crew, and carry accountability for the outcome. Copilot's pushback is sophisticated, but it is still pattern-matched

critique rather than genuine adversarial reasoning. It may miss highly specific or contextual weaknesses — the ones that come from having worked this exact hall before — that a real expert would catch. Use it as a first filter, not the only one.

The prompting template:

Play **the** role of [specific skeptical counterpart].
I am going **to** present [my idea/proposal/analysis].
Push **back** hard. Don't flatter **me**. Identify **every** weakness,
every assumption, **every** place **where** my reasoning **is**
vulnerable. Be direct.

[Present your idea/proposal/analysis]

GES examples:

Play **the** role of a skeptical procurement lead **at** a show organizer **that** runs a \ \$200M annual event portfolio **and** has been told **to** cut supplier costs **by** 12%. I'm going **to** pitch you **on** renewing GES **as** your general service contractor **with** an expanded scope **that** includes exhibitor success outreach **and** sponsorship consulting. Push **back** hard.

Don't accept **my** assumptions. Tell **me** what objections you'd actually raise.

Here **is** my pitch: [present your pitch]

Play **the** role of a venue safety officer conducting **a** rigging **and** overhead hang review. I am going **to** present our structural plan **for** **a** 60x80 island stand **with** **a** suspended banner ring **and** integrated LED. Identify every area where you would ask hard questions, where our documentation is thin, **and** where our **load** calculations may **not** hold up under scrutiny.

Here is **the** plan: [paste **the** plan]

Play the **role of** an exhibiting brand's marketing director who has just been handed our quote and thinks material handling is a made-up line item designed to inflate the bill. I'm presenting a \ \$7.5M program for a multi-city custom exhibit build with full door-to-floor logistics. Find the weaknesses in my **cost** story. **Do not** soften your feedback.

Here **is** my proposal **summary**: [paste the **summary**]

Play the role of an operations director who has watched four international shows blow their **move-in** window in the last two years. I'm proposing a compressed two-day **move-in** for a 900-booth show in a venue with six docks **and** a single freight elevator to the mezzanine. Tell me exactly where this plan breaks.

Here is my sequencing plan: [paste the plan]

The Sparring Partner and Overconfidence

The Sparring Partner technique is extraordinarily useful — but it carries one risk worth naming. Copilot's pushback is sophisticated enough that it may feel comprehensive. It may not be. There are venue-specific realities, union jurisdiction nuances, and relationship factors that a human expert will surface that Copilot will miss entirely — because they were never written down anywhere Copilot can reach.

Run your ideas through the Sparring Partner to improve them. Then also run them through an actual human colleague you trust — ideally one who has worked that building. The combination — AI critique followed by human review — produces the best outcome.

*And to be unambiguous: **never** treat an AI-role-played safety review as a substitute for a real one. Safety sign-offs come from qualified humans. Full stop.*

"Provide excellent service and execution."


The Sparring Partner technique is Excellence operationalized. Excellence is not the absence of problems; it's finding them earlier than anyone else does.

There is a specific professional maturity in deliberately inviting criticism of your own work before someone else delivers it for free, at a worse moment, in front of a client. The people at GES who build a habit of adversarial self-review — who spar with their own plan on Wednesday so the pre-con call on Friday is boring — are the ones whose shows open clean.

Excellence at scale means doing that 4,000 times a year. That's only possible if the pressure-test is cheap. Copilot makes it cheap.

6. The Goldilocks Zone of Prompt Length

Now that you understand the five techniques, there is one meta-principle that governs all of them: **the length and precision of your prompt determines the quality of your output, and there is a precision sweet spot — not too short, not too long.**



The "too short" failure mode:

"Write a memo about our Q2 show operations performance."

This prompt has no context (who is reading it? what do they already know?), no specific goal (is this for leadership? the organizer? the regional GM?), no format guidance (length? tone? structure?), and no constraints (what should it include? what should it avoid?). Copilot will generate something. It will be polished. It will be almost entirely useless for your specific purpose, because it is optimized for the generic version of the task rather than your version.

The "too long" failure mode:

Some users, having learned that detail helps, overcompensate. They write prompts that are five paragraphs of background, instruction, counter-instruction, caveats, and competing requirements. At a certain level of complexity, this actually degrades output — the model gets pulled in multiple directions by conflicting guidance, and the result is incoherent or over-hedged.

The Goldilocks zone — the four-component prompt:

The research and practitioner experience on prompting converges on a consistent structure for most business tasks. A good prompt has four components:

1. **Context:** Who you are, what situation you're in, what Copilot needs to

know to frame its response appropriately

2. **Goal:** What you want Copilot to produce — specific, not vague
3. **Format:** How you want the output structured (bullet list, formal memo, executive summary, table, numbered recommendations, etc.)
4. **Constraints:** What to include, what to avoid, tone, length, any requirements that bound the output

These four components can usually be covered in three to five sentences, or a short paragraph. That is the sweet spot.

The GES four-component prompt in action:

Context: I'm an account manager at GES preparing for a renewal conversation **with** a show organizer whose annual **event** represents roughly **\\$15M in** services across three halls. They have been a GES client **for** eight years. Execution has been strong, but a competing general service contractor has approached them **with** a more aggressive material handling rate.

Goal: Draft talking points **for** the renewal conversation that emphasize our partnership history, **global** footprint, **and** execution record without immediately conceding **on** rate.

Format: A bulleted list **of** five **to** seven talking points, **each with** a one-sentence expansion.

Constraints: Keep the tone confident but **not** defensive. **Do not** include any specific rate numbers **or** commitments – I'll handle those **in** the room. The organizer's team is highly experienced and has heard every standard GSC pitch before.

That prompt takes thirty seconds to write. The output it produces is immediately usable — not a generic draft that requires fifteen minutes of editing to make relevant.

Here is a second one from a completely different part of the business, to show the structure travels:

Context: I'm a logistics coordinator managing inbound freight **for** a 1,200-booth show **at** a venue **with** eight docks **and** a three-day move-in. About 35% of exhibitor freight **is** routing **through the** advance warehouse; **the rest is** direct-to-show. We had significant marshaling yard congestion **at** this venue **last year**.

Goal: Produce a freight target schedule narrative I can send **to the** show team explaining **the** sequencing logic **and the** three highest-risk windows.

Format: A short memo – one **paragraph of** context, a table **of** target windows **by** hall, **then** a numbered **list of** risks **with** a mitigation **for** each.

Constraints: Assume standard union jurisdiction **for** this venue. Do **not** invent specific truck counts; use placeholders **where** I need **to** fill **in** actuals. Keep **it** under one page.

flowchart LR

A["Context\n(Who + Situation)"] --> E["The Precision\nSweet Spot"]

B["Goal\n(Specific Output)"] --> E

C["Format\n(Structure)"] --> E

D["Constraints\n(Boundaries)"] --> E

E --> F["Professional-Grade\nCopilot Output"]

style E fill:#1a73e8,color:#fff

style F fill:#f4a400,color:#fff

A note on iterating:

You rarely need to get the perfect prompt on the first try. Copilot conversations are threaded — your follow-up messages have context from everything that came before. Think of prompting as a conversation: start with a good four-component prompt, review the output, and then refine with specific follow-up instructions. *"That's good — now make the tone more direct"* or *"Remove the second and fourth points and expand the third into two separate bullets"* or *"Rewrite this for a first-time exhibitor who doesn't know what drayage means."*

That last one is worth dwelling on. The exhibition industry runs on jargon — drayage, targets, GSC, EAC, I&D, carnet — and roughly 150,000 exhibitors a year encounter that jargon for the first time. Asking Copilot to translate an internal document into plain language for a first-time exhibitor is one of the highest-value, lowest-effort prompts available to anyone in exhibitor services. The follow-up prompt is often where the real refinement happens.

7. From Prompt to Assignment: Prompting for Copilot Cowork

Everything up to this point has been about **Copilot Chat prompting** — the back-and-forth conversation where you steer each step. That skill is foundational and it is not going away.

But there is a second mode, and it demands a different kind of prompt.

Microsoft Copilot Cowork — announced in the Frontier early-access program in March 2026 and generally available worldwide since **June 16, 2026** — executes complex, long-running, multi-tool tasks end-to-end across Microsoft 365 and returns **finished artifacts**. Not a draft in a chat window: a deck saved to SharePoint, a workbook with labeled tabs, an email queued for your approval, a meeting on the calendar. At GA, more than half of the Fortune 500 had adopted it, and Microsoft reported it as the fastest-growing feature in the history of the Frontier program.

The crucial property for GES: **Cowork runs in a protected cloud environment, which means tasks keep running even when your laptop is off**. A show manager can assign a post-show reconciliation package from the floor in Las Vegas at 7:00 a.m., work an eleven-hour move-out, and review a finished draft on their phone that evening. That is not a footnote. For a company running 4,000+ events a year across 75+ countries with lean teams and immovable show dates, it is close to the whole point.

Chat vs. Cowork vs. Agents:

Choosing Your Mode

Dimension	
Conversational	Ready-made

Best for	AI for drafting, Q&A, and ideation	Delegate and execute long-running, multi-step work across your apps	agents for specialized or repeatable tasks
How you interact	A conversation: you steer each step from prompt → response	An assignment: you describe the goal, check in at key milestones	A workflow: you pick an agent built for a specific job
Typical work pattern	You're in the loop — one prompt, one result, then you choose what's next	You step away — Cowork plans, manages files and tasks across apps, and delivers completed work	You run it on demand — the agent handles the same scoped task each time
GES example	"Rewrite this exhibitor notice in plainer language."	"Build the full post-show reconciliation package for the Chicago show."	"Triage inbound exhibitor service kit questions against the standing scope."

Microsoft's Five-Part Prompt Structure for Cowork

The four-component prompt (Context, Goal, Format, Constraints) is the right structure for Chat. Cowork needs one more thing — because Cowork **takes actions**, and actions need boundaries.

Microsoft's official guidance defines a **five-part structure**:

The Five-Part Cowork Prompt Structure

Part	What It Means	GES Example
Outcome	One sentence describing what <i>done</i> looks like	"A complete post-show reconciliation package for the Chicago show, ready for the organizer's finance team."
Inputs	The specific people, files, sites, or time ranges the	"The show SharePoint library, the final labor timesheets, the freight manifest workbook, all change orders logged

	task should use	between move-in and move-out, and the pre-con call notes."
Definition of done	The concrete deliverable — a document saved, an email sent, a meeting booked	"An Excel workbook with tabs for labor, freight, and change orders; a two-page Word summary; both saved to the show's SharePoint folder."
Constraints	Things to avoid or honor	"Do not contact the organizer. Use the FY27 reconciliation template. Flag any variance over 5% rather than resolving it. Keep the summary under two pages."
Approval scope	Which actions you want to review explicitly, beyond the default checkpoints	"Ask me before creating any new SharePoint folders. Ask me before any email — internal or external."

Before and after — the difference this makes:

Before (too open):

Help me **close out** the Chicago **show**.

Cowork has to guess at nearly everything: which show, which files, what "close out" means, what a finished product looks like, and whether it's allowed to email anyone. Guessing shapes results, and not in your favor.

After (five-part):

Outcome: A complete post-show reconciliation package for the Chicago show that our account manager can walk the organizer's finance team through next Tuesday.

Inputs: The Chicago show SharePoint library, the **final** labor timesheets, the freight **and** material handling manifest workbook, every change order logged between **move-in** and move-out, **and** my notes from the post-show debrief in Teams.

Definition of done: One Excel workbook with separate labeled tabs for labor, freight/drayage, **and** change orders, plus a two-page Word summary memo. Both saved to the Chicago show folder in SharePoint.

Constraints: Use the FY27 reconciliation template. Do **not** contact the organizer **or** any exhibitor. Flag any line-item variance over 5% in a clearly marked exceptions section rather than trying to resolve it. Keep the memo under two pages. Use GES terminology consistently – "material handling," **not** "drayage fees."

Approval scope: Ask me before creating any **new** folders, **and** ask me before sending any email, internal **or** external.

Three more GES Cowork assignments worth stealing:



{tab-item} Show Ops — Pre-Con Packet

Outcome: A complete pre-con briefing packet **for the** show team ahead **of the** Tuesday planning call.

Inputs: The current floor plan revision, **the** freight target schedule,
the labor forecast workbook, **the** venue site survey writeup, **and the** last **six** weeks **of** email **with the** organizer.

Definition **of** done: A PowerPoint deck (**12** slides **max**) plus **a one-** page

Word summary **of open items with** owners, saved **to the** show SharePoint **folder, and a** calendar hold **for a 30-**minute internal prep session **the day before.**

Constraints: Do **not include** exhibitor-level pricing. Use **the** standard

GES pre-con deck template. Any **date or** weight you cannot verify **from the** source **files** should be marked "**TBC,**" **not** estimated.

Approval scope: Ask **before** booking **the** calendar hold. Ask **before** posting anything **to the** show Teams channel.

*Outcome: A renewal briefing pack **for the** account team ahead **of the** organizer's procurement review.*

*Inputs: Three years **of** post-show reconciliations **for** this organizer, **the** current scope **of** work, satisfaction survey results, **and** all correspondence **with the** organizer's procurement lead this year.*

*Definition **of** done: A Word briefing document **with** a performance summary, a risk section, **and** five anticipated objections **with** suggested responses; plus an Excel **tab** showing **year-over-year** service volume **and** variance.*

*Constraints: No rate proposals **or** pricing commitments. Cite **the** source **file for every number**. Keep **the** briefing under six pages.*

*Approval scope: Ask **before** sending anything **to** anyone. This **is** internal-only **until I say** otherwise.*

*Outcome: A quarterly reuse **and** materials summary supporting our Better Stands reporting **for the** EMEA region.*

*Inputs: The production facility inventory workbooks, **the** show-by-show build records **for the** quarter, **and the** reuse tracking sheet.*

*Definition of done: An Excel workbook **with a** tab per facility **and a** consolidated summary **tab**, plus **a** one-page narrative highlighting **the three** largest reuse gains **and the three** largest gaps.*

*Constraints: Use only data present **in the** source **files** – **do not** estimate **or** extrapolate **any** emissions figure. Flag missing data explicitly **in a** "data gaps" section.*

*Approval scope: Ask **before** creating **any new** SharePoint location.
No emails.*

Reviewing Like a Manager

Here is the genuine behavioral shift, and it is worth naming plainly.

The prompt-craft skill of Copilot Chat is describing a task. The skill of Cowork is describing an outcome and then reviewing like a manager.

That is a different professional muscle. It is the difference between doing the install yourself and running the crew. Cowork will pause and ask permission before sensitive actions — sending an email, posting in Teams, updating a record. You will see a rich preview for most of these, a risk-level indicator for medium and high risk actions, and options to approve once, approve for the rest of the session, scope approval to a specific recipient or domain, approve everything pending at once, or cancel.

Microsoft's own guidance is blunt about this: **always review details before approving — check recipients, content, and other details.** Every task runs with your permissions and sees only what you can see. Data stays in your tenant. Actions are auditable. And people — not Cowork — remain responsible for business decisions.

Cowork's approval prompts are the exact point where two GES values meet the software.

Trust is what makes delegation possible at all — trust in the permission architecture, trust in your colleagues, trust that the audit trail exists.

Responsibility is what makes delegation safe: the approval button has your name on it.

*The failure mode to watch for is approval fatigue. Clicking **Approve All (7)** because you're between halls and the badge line is backing up is exactly how an unfinished draft reaches an organizer. Scope your approvals deliberately at the start of a task — that's what the approval scope line in your prompt is for — and then actually read the previews for anything that leaves the building.*

An email you didn't read but did send is still an email you sent.

8. Try This: Run the Same Show Floor Question Five Ways

Here is where the theory becomes practice. This exercise produces one of the clearest demonstrations of how much the technique matters — more than the question itself.

The baseline question:

"How should GES approach a long-standing show organizer whose event is shrinking — fewer exhibitors, less square footage, and pressure to cut service costs — at the same time their contract comes up for renewal?"

This is a real strategic question that any account manager, show manager, or regional GM at GES might face. It is specific enough to be meaningful, general enough to work without confidential client data.

Run it five times — once with each technique:

{tab-item} Technique 1: Role-Based

Act **as a** veteran regional general manager **at a global** general service contractor **with 15** years **of** experience managing declining **and** turnaround show accounts. How should GES approach **a long-standing** show organizer whose event is shrinking — fewer exhibitors, less square footage, **and** pressure **to** cut service costs — **at the same time** their contract comes up **for** renewal?

What to notice: The response should have a notably more specific, risk-focused, and action-oriented character than a generic answer. It should reference concrete levers — scope restructuring, floor plan redesign to make a smaller show feel full, sponsorship revenue consulting, exhibitor success programs to slow attrition — rather than abstract "partnership" language.

Walk me through your reasoning step by step – considering the organizer's economics, our cost base, the competitive landscape,
and the **long-term** value of the account – before giving me a **final** recommendation on how GES should approach a **long-standing** show organizer whose event is shrinking at the same time their contract comes up for renewal.

What to notice: The response should show its work — market context, organizer's revenue position, our margin structure on this account, switching costs, competitive threat assessment, option evaluation — before landing on recommendations. The reasoning chain is the value.

Here are **two** examples **of** how our account teams have handled similar situations well:

Example **1** – Strong: [paste **a** real **or** hypothetical example **of** a well-handled renewal **with** **a** contracting show]

Example **2** – Strong: [paste another strong example]

Now give me **a** recommendation **for** how GES should approach **a** **long-standing** show organizer whose event is shrinking **at the** same **time** their contract comes up **for** renewal. Match **the** style **and** specificity **of the** strong examples.

What to notice: The output should mirror the structure and depth of your examples. If your examples were concise and action-oriented, the new output will be too.

*I need a recommendation on how GES should approach a **long-standing***

show organizer whose event is shrinking – fewer exhibitors, less

*square footage, **and** pressure to cut service costs – at the same*

time their contract comes up for renewal. Before you give me an

answer, ask me any questions you need to give me the most useful

possible recommendation.

What to notice: Copilot should ask about the size and margin of the account, how long the decline has run, whether the decline is industry-wide or organizer-specific, which competitors are circling, what scope we currently hold, and how much of the GES Collective (Spiro, onPeak, Visit, SHOWTECH) is already engaged. Your answers will produce a significantly more tailored recommendation.

Play **the** role **of** a skeptical regional finance director who has watched too many account teams protect revenue **by** quietly giving away margin. I'm going **to** propose **that** we renew this shrinking show **at** a reduced material handling rate **in** exchange **for** a three-**year** term **and** expanded scope across onPeak housing **and** Visit registration. Push **back** hard. Tell **me** **where** this approach **is** risky **and** what I'm **not** thinking **about**.

My proposal: We hold **list** pricing **on** labor, concede **8%** **on** material handling, **and** bundle onPeak **and** Visit **at** a blended discount, **in** exchange **for** a three-**year** exclusive **and** a joint exhibitor growth program **with** shared targets.

What to notice: This output should be genuinely challenging. A good Sparring Partner response will identify risks like margin erosion that persists after the show recovers, the precedent set with other organizers who will hear about the rate, whether "expanded scope" actually carries incremental margin or just incremental cost, the enforceability of shared growth targets, and what happens in year three if the show keeps contracting.

Compare the five outputs. They should be meaningfully different — not just in tone, but in substance, depth, and practical utility. That difference is what you're learning to produce deliberately.

Bonus round — run it a sixth way, as a Cowork assignment. Rewrite the same question using the five-part structure from Section 7: Outcome, Inputs, Definition of done, Constraints, Approval scope. Notice how differently you have to think. Chat asks "what do I want to talk about?" Cowork asks "what do I want to exist when I come back?"

What to Record After the Exercise

Keep notes on:

- *Which technique produced the most immediately useful output for this type of question?*
- *Which technique surprised you most?*
- *Where would you combine techniques? (Role-Based + Chain-of-Thought is a particularly powerful combination for complex analytical tasks.)*
- *What would you add to refine each prompt further?*
- *What changed when you restructured the request as a Cowork assignment?*

Your prompting instincts improve with every iteration. The professionals who become power users of Microsoft 365 Copilot are not the ones who got lucky on the first try — they are the ones who treated each prompt like a hypothesis to be tested and refined.

9. Putting It Together: Your First Week Prompting Plan

This course gives you the frameworks. What moves you from "I understand the techniques" to "I use them automatically" is repetition — applied to real GES work, on real shows, this week.

Here is a concrete five-day prompting plan:

GES First-Week Prompting Plan

Day	Technique	Task to Try at GES
Monday	Role-Based Prompting	Assign a relevant professional role — show operations director, venue safety officer, skeptical organizer — and use it to review a document you're currently working on

Tuesday	Chain-of-Thought	Use Copilot to reason through a decision you need to make this week — a freight routing call, a labor staffing level, a scope question — asking for step-by-step reasoning before the conclusion
Wednesday	Few-Shot	Pull two or three examples of a document type your team produces regularly (site survey, reconciliation, design spec) and ask Copilot to generate a new one in the same style
Thursday	Reverse Prompting	Let Copilot interview you for an organizer presentation or proposal you need to develop — answer its questions, then review the output
Friday	Sparring Partner	Present one of your current plans or recommendations to Copilot in Sparring Partner mode and see what comes back

By Friday, you will have hands-on experience with all five techniques applied to real work. That is more practical prompting practice than most Copilot users accumulate in their first three months.

Week two, if you want to keep going: take the single task from week one that produced the most useful output and rewrite it as a five-part Cowork assignment. Then walk away and see what comes back. That is the transition from user to delegator — and it is where the compounding really starts.

10. Productive Struggle Problem

You are a senior account manager at GES. Your most important organizer client — an association whose annual exposition fills three halls, draws 1,400 exhibitors, and represents roughly \$22M in annual services across the GES Collective — has just informed you that a competing general service contractor has proposed taking over the full account, claiming it can save the organizer \$180,000 annually on material handling and labor alone.

You have thirty minutes before the organizer's executive director calls to discuss. You need: a clear understanding of whether the competitor's savings claim is plausible, a retention strategy that does not simply concede on rate,

and talking points that draw on the eight-year relationship history and the parts of the GES Collective — Spiro, onPeak, SHOWTECH, Visit by GES — that a single-service competitor cannot match.

You have access to Microsoft 365 Copilot and the five techniques from this chapter.

The challenge: Design a sequence of Copilot prompts — using at least three of the five techniques — that you would actually run in the next twenty-five minutes to prepare for this call. Write out each prompt in full, explain which technique it uses and why you chose it for that step, and describe what you expect from each output.

Extension (optional but recommended): Assume you had five days' notice instead of thirty minutes. Write the same preparation as a single **Cowork assignment** using the five-part structure — Outcome, Inputs, Definition of done, Constraints, Approval scope. Then compare: what does the extra time and the delegation model let you produce that thirty minutes of Chat prompting cannot?

There is no single right answer. There are better and worse sequences, and the quality of your thinking about *why* you're choosing each technique is as important as the prompt itself.

Discussion

Prompting as a Professional Skill

The five techniques in this chapter are not software features — they are communication frameworks that work because of how language models process context. As Microsoft continues to evolve Microsoft 365 Copilot, the specific mechanics will change. The principles will not.

Consider: In what ways does prompting Copilot resemble briefing a highly capable new coordinator on their first show? In what ways does it differ? What does that comparison reveal about where human judgment — the kind earned by standing in a hall at 4:00 a.m. — remains irreplaceable in the AI-assisted workflow?

Discussion Guidelines:

Your response should engage substantively with the comparison between prompting AI and briefing a human colleague. Include at least one specific example from your own professional experience — either an experience where clear communication dramatically improved a colleague's output, or a case where a vague request produced work that missed the mark entirely. If you have worked a show floor, use it: the gap between "set the booth" and a real set of I&D instructions is the same gap this chapter is about. Apply that experience to what you've learned about prompting.

Support your perspective with at least one credible source — this might be something from the Microsoft 365 Copilot official documentation, the Skilling Center resources at adoption.microsoft.com/copilot/skilling-center/, or a relevant piece of research on human-AI collaboration in professional settings.

After posting your response, engage with **at least two classmates** by extending or challenging a specific claim they made — not just affirming it. "I agree because..." is not a sufficient peer response. "I'd push back on your claim that X because I've seen Y, which suggests Z" is.

Do not summarize or repeat what you read. Share what you think — grounded in evidence — about what this means for how GES professionals should work going forward.

Glossary

Microsoft 365 Copilot

The paid, enterprise-grade AI assistant available through a Microsoft 365 Copilot license. Unlike Copilot Chat, M365 Copilot connects to your organizational data via Microsoft Graph and operates within your M365 security boundary.

Copilot Chat

The free, general-purpose AI assistant available at copilot.microsoft.com. Provides web-connected AI assistance but does not access organizational data (emails, files, Teams chats) by default.

Microsoft Graph

The intelligence layer within Microsoft 365 that maps organizational data — emails, meetings, chats, calendar events, files, and shared

documents. Copilot uses Microsoft Graph to ground its responses in your actual work context.

Grounding

The process by which Microsoft 365 Copilot accesses Microsoft Graph to retrieve relevant organizational context before generating a response. Grounding is what makes Copilot responses personalized to your actual work rather than generic.

Copilot Agent

A purpose-built AI assistant configured for a specific workflow, knowledge base, or organizational function. More specialized than the general M365 Copilot experience; configured through Microsoft Copilot Studio.

Copilot Cowork

The Microsoft 365 mode that executes complex, long-running, multi-step tasks across apps and returns finished artifacts rather than chat responses. Generally available worldwide since June 16, 2026. Tasks continue running in the cloud even when your device is off.

Five-Part Prompt Structure

Microsoft's official prompting structure for Copilot Cowork: Outcome, Inputs, Definition of done, Constraints, and Approval scope. Extends the four-component Chat prompt by adding explicit boundaries on the actions Cowork is permitted to take.

Approval Scope

The part of a Cowork prompt that specifies which actions you want to review explicitly before they execute — beyond Cowork's default checkpoints for sensitive actions such as sending email or posting in Teams.

m365.cloud.microsoft

The unified web portal for Microsoft 365 Copilot — the "front door" that provides access to all M365 Copilot capabilities, the Prompt Gallery, and connected M365 apps from a single authenticated session.

Microsoft 365 Copilot Search

A universal search capability within the M365 Copilot experience that searches across all connected M365 apps (Outlook, SharePoint, Teams, OneDrive) and third-party data in a single query.

Role-Based Prompting

A prompting technique that assigns a specific professional identity or perspective to Copilot before making a request, shifting the analytical frame, vocabulary, and focus of the output.

Chain-of-Thought Reasoning

A prompting technique that asks Copilot to show its reasoning step-by-step before giving a final answer, producing more transparent, coherent, and auditable analytical output.

Few-Shot Prompting

A prompting technique that provides Copilot with a small set of high-quality examples before asking it to produce new output — teaching by demonstration rather than by description.

Reverse Prompting

A prompting technique that inverts the standard dynamic by asking Copilot to interview you with clarifying questions before producing output, ensuring the request is well-understood before work begins.

Sparring Partner

A prompting technique that assigns Copilot the role of a skeptical, critical counterpart who pushes back on your ideas — used to pressure-test thinking before high-stakes conversations such as an organizer procurement review.

Goldilocks Zone

The optimal prompt length and specificity for a given task — precise enough to produce relevant output, concise enough to avoid confusing or overloading the model. Typically achieved with the four-component prompt structure: Context, Goal, Format, Constraints.

Four-Component Prompt

A prompting structure with four elements — Context, Goal, Format, and Constraints — that consistently produces more useful output than vague or overly complex prompts.

Prompt Gallery

Microsoft's official curated library of tested, verified prompts organized by job function and application. Available at m365.cloud.microsoft/copilot-prompts.

Copilot Skilling Center

Microsoft's official learning hub for Microsoft 365 Copilot, with structured learning paths and role-specific guidance. Available at adoption.microsoft.com/copilot/skilling-center/.

Permission Scoping

The architectural principle that Microsoft 365 Copilot only surfaces information the signed-in user already has permission to access — ensuring Copilot cannot bypass existing security and compliance controls.

GSC (General Service Contractor)

The official on-site services provider appointed by a show organizer to deliver decorating, material handling, labor, signage, and floor operations for an event. GES's core role at most of the 4,000+ events it produces each year.

EAC (Exhibitor Appointed Contractor)

A third-party contractor engaged directly by an exhibitor rather than by the show organizer. Spiro frequently operates in this capacity for custom brand experiences.

Material Handling / Drayage

The receipt, transport, storage, and delivery of exhibitor freight between the dock or advance warehouse and the booth space on the show floor, plus the reverse at move-out.

Freight Target

A scheduled arrival window assigned to inbound exhibitor freight to sequence dock access and prevent marshaling yard congestion during move-in.

Advance Warehouse

A facility that receives exhibitor freight in the weeks before a show, stores it, and delivers it to the booth on schedule — reducing the risk and congestion of direct-to-show shipping.

Marshaling Yard

The staging area where carriers check in and wait to be called to a dock during move-in and move-out.

I&D (Install & Dismantle)

The labor operation of building exhibits before a show opens and tearing them down after it closes.

Move-In / Move-Out

The contracted windows during which exhibits are installed before the show and removed after it — the two highest-pressure periods in the event lifecycle.

Exhibitor Service Kit

The document package issued to exhibitors containing ordering deadlines, rates, shipping instructions, venue rules, and service options for a specific show.

Carnet

An ATA Carnet — an international customs document permitting temporary, duty-free import of exhibition goods across borders. Essential for GES's work across 75+ countries.

Union Jurisdiction

The rules governing which trade or labor group may perform which tasks in a given venue or region — a primary driver of labor planning and cost in North American exhibition work.

Site Survey

A structured on-site assessment of a venue capturing dock configuration, freight elevator capacity, ceiling heights and rigging points, egress requirements, jurisdiction notes, and known operational constraints.

Post-Show Reconciliation

The closeout process that reconciles ordered versus delivered services, forecast versus actual labor, estimated versus actual freight, and all change orders — producing the final billing and performance record for a show.

The GES Collective

The GES brand portfolio: GES Exhibitions (core GSC services), Spiro (global experiential agency and EAC), onPeak (event accommodations and housing), SHOWTECH (power and lighting), and Visit by GES (event management technology — registration, lead capture, visitor apps, NFC Touchpoints).

T.R.U.E. Values

The GES core values: Trust, Responsibility, Understanding, and Excellence. Applied to AI work, they translate to honesty about how a document was produced, ownership of every number you ship, empathy for the human who receives your output, and the discipline to find problems before anyone else does.

Leader's Takeaway

Microsoft 365 Copilot is not a feature you turn on. It is a capability you develop — in yourself, and in your team.

The infrastructure is in place. The security architecture is sound. The intelligence layer that grounds Copilot in your organizational context is active. The tools are available at m365.cloud.microsoft right now, and Cowork has been generally available since June 2026.

What remains is the skill — specifically, the prompting skill that determines whether Copilot becomes a transformative part of how GES works or sits as an underutilized line item in the software budget.

The five techniques in this chapter are not exotic. They are communication principles applied to a new kind of tool. Role-Based Prompting is how you route a question to the right specialist. Chain-of-Thought is how you ask for the cost breakdown, not just the quote. Few-Shot is how you onboard a new coordinator with examples of your house standard. Reverse Prompting is how you let the expert run discovery before the work begins. The Sparring Partner is how you walk the floor plan with your harshest critic on Wednesday so the pre-con call on Friday is uneventful.

There is a reason this matters right now specifically. On **December 31, 2024**, GES completed its separation from Viad and became an independent company for the first time in 55 years. That independence means the roadmap belongs to GES — the investment decisions, the technology choices, the pace of change. Companies do not get many moments like that. What you do with the first few years of it tends to define the next decade.

The onPeak AI Smart Suite already proves the organization can build and ship AI that works. The question this chapter answers is whether that capability stays concentrated in one product team or spreads to every account manager, show operations lead, exhibit designer, freight coordinator, and exhibitor services rep across 24 facilities and 75+ countries.

These techniques are learnable. They improve with repetition. And the colleagues who develop fluency with them fastest will not just be more productive individually — they will become the informal teachers who spread that capability to everyone around them. That is how a company of 2,600 people gets good at something quickly: not through a mandate, but through the person two desks over who found a better prompt and shared it.

That is the compound return on investing in this skill. It is not linear — it multiplies. Across 4,000 shows a year, it multiplies a lot.

Continue Building

- ***Chapter 4*** established the mindset — where AI helps, where it doesn't, and why judgment stays human.
- ***Chapters 6-12*** apply these five techniques inside Word, Excel, PowerPoint, Outlook, Teams, SharePoint, and OneNote, using real GES artifacts: exhibitor service kits, labor forecasts, organizer decks, show document libraries, and site survey notebooks.
- ***Chapter 14*** goes deep on Copilot Cowork — the five-part prompt structure introduced here, the approvals model, scheduled and event-driven tasks, and the governance questions that come with delegating real work.
- ***Microsoft's Copilot Prompt Gallery*** — m365.cloud.microsoft/copilot-prompts
(<https://m365.cloud.microsoft/copilot-prompts/>)
- ***Microsoft 365 Copilot Skilling Center*** — adoption.microsoft.com/copilot/skilling-center
(<https://adoption.microsoft.com/copilot/skilling-center/>)